

AWS Auto Scaling & CI/CD for a 3-Tier Application

AWS Auto Scaling & CI/CD for a 3-Tier Application

Assignment Book

Student	Joyanta Sarker Joy
Course	Mastering DevOps — Batch 13
Module	9 — AWS Auto Scaling & CI/CD for a 3-Tier Application
AWS Region	Asia Pacific (Mumbai) — <code>ap-south-1</code>
Repository	https://github.com/JsJoyBitopibd/-AWS-Auto-Scaling-CI-CD-for-3-Tier-Application-Part-1-Auto-Scaling-for-3-Tier-Applications
Date	25 - 27 September 2026

Contents

1. Overview and account constraints
 2. Part 1 — Auto Scaling for a 3-Tier Application
 3. How the servers talk to each other (with commands)
 4. Part 2 — CI/CD Pipeline
 5. Appendix — scripts
-

1. Overview

Student	Joyanta Sarker Joy
Course	Mastering DevOps — Batch 13
Module	9 — AWS Auto Scaling & CI/CD
AWS Region	Asia Pacific (Mumbai) — <code>ap-south-1</code>

What this project does

Part 1 converts a backend that ran on a single EC2 instance into a **highly available, auto-scaling, self-healing** 3-tier architecture: an Application Load Balancer in front, a fleet of backend servers managed by an Auto Scaling Group across two Availability Zones, and a MySQL database tier behind a locked-down firewall.

Part 2 adds an automated **CI/CD pipeline** (Source → Build → Test → Deploy) that ships new versions of the backend to the auto-scaling fleet, proves that freshly launched servers configure themselves automatically, and finishes with a **Blue/Green deployment** plus a tested **rollback**.


How Part 2 was built without IAM roles

The account cannot create IAM roles, which CodePipeline, CodeBuild and CodeDeploy all require. The pipeline therefore uses **GitHub Actions** (Source → Build → Test → Package → Release) and a **pull-based deployer** on every server (`scripts/deploy.sh` , run by a `systemd` timer every 60 s and at first boot). A `git push` reaches the whole fleet in under two minutes with no AWS credentials anywhere; new Auto Scaling instances fetch the current release before their first health check. Blue/Green is a second target group + Auto Scaling Group pinned to a *candidate* release, switched by the ALB listener. Details and evidence: [Part 2](#).

Architecture

Part 1 — Highly-Available 3-Tier Auto Scaling Architecture

AWS Region: Asia Pacific (Mumbai) · ap-south-1 | Joyanta Sarker Joy · Mastering DevOps (Batch 13)



Request flow: ① **Browser → ALB (:80) → ② ALB → healthy backend EC2 (:3000, health-checked on /health) → ③ Backend → MySQL (:3306).**

Documentation (repository layout)

Document	What it covers
Part 1 — Auto Scaling for 3-Tier Applications	VPC, security groups, database tier, Launch Template, Target Group, ALB, Auto Scaling Group, health checks, self-healing test, scaling policy + load test — with screenshots
Part 2 — CI/CD Pipeline	Pipeline architecture, source/build/deploy stages, auto-deploy to new instances, Blue/Green deployment, rollback test — with screenshots
Assignment Book (PDF)	Everything below in one document for submission — cover, both parts, the communication guide, all screenshots
How the servers talk to each other	Every network hop in the system explained in plain words, with the real commands used at each hop
Terraform audit	Read-only Terraform program that reports which of the hand-built resources still exist (PRESENT / MISSING) and prints the values needed to continue — how to run it on Windows, step by step

Repository structure

```

├── README.md                ← you are here
├── docs/                    ← the full write-ups (assignment documentation)
│   ├── part1-auto-scaling.md
│   ├── part2-cicd.md
│   └── how-servers-communicate.md
├── diagrams/                ← architecture diagrams (PNG + editable SVG)
├── SS/                      ← evidence screenshots (P1-xx = Part 1, P2-xx = Part 2)
├── app/                     ← the Node.js backend application
│   ├── app.js               ← serves / /health /version , reads from MySQL
│   ├── package.json         ← version here = release tag
│   └── test/health.test.js   ← tests run by the pipeline
├── scripts/                 ← EC2 user-data (first-boot) scripts
│   ├── db-server-userdata.sh ← builds the MySQL database server
│   ├── app-server-userdata.sh ← Launch Template v1 (Part 1: app embedded)
│   ├── app-server-userdata-v2-blue.sh ← Launch Template v2 (Part 2: pull-based deployer,
BLUE)
│   ├── app-server-userdata-v3-green.sh ← Launch Template v3 (Blue/Green: GREEN, pinned
candidate)
│   ├── deploy.sh            ← the deployer that runs on every server
│   └── fix-ssh-key-permissions-windows.ps1 ← makes a .pem key usable by Windows OpenSSH
├── terraform/               ← read-only audit: is every resource still there? (main.tf,
audit.tf, README.md)
└── .github/workflows/ci-cd.yml ← the CI/CD pipeline (GitHub Actions)

```

Account constraints and how the design adapts

This project was built in a **restricted course account**. The limits below are real, and the design works around every one of them:

Restriction	Impact	How we adapted
Cannot create IAM roles	Native CodeDeploy/CodeBuild/CodePipeline need service roles + an EC2 instance profile	Part 1 needs no user-created roles at all. Part 2 uses existing roles if present, otherwise a role-free pipeline (see Part 2 doc)
No SSM permission	Can't use Session Manager or Parameter Store	SSH with a key pair for access; environment variables for app config
Cannot create RDS (<code>rds:CreateDBSubnetGroup</code> denied)	No managed database	MySQL (MariaDB) self-hosted on a small EC2, auto-configured by user-data, protected by <code>bitopi-db-sg</code>
No CloudShell	Can't run AWS CLI in the console	All configuration done through the console; permissions confirmed by building

Note on secrets: the database password in the scripts is a throw-away lab credential and every resource is destroyed at the end of the assignment. In production this would live in

Progress

- [x] Part 1 architecture diagram
 - [x] VPC + subnets across 2 AZs
 - [x] Security groups (ALB / Backend / DB)
 - [x] Database tier (MySQL on EC2)
 - [x] Launch Template
 - [x] Target Group + Application Load Balancer
 - [x] Auto Scaling Group
 - [x] Health checks + self-healing test
 - [x] Scaling policy + load test (scale-out / scale-in)
 - [x] Part 2 CI/CD architecture diagram
 - [x] CI/CD pipeline (Source → Build → Deploy)
 - [x] Auto-deploy to new Auto Scaling instances
 - [x] Terraform audit — read-only check that every resource is still present (`terraform/`)
 - [x] Blue/Green deployment (green ASG on LT v3, tested on a :8080 test listener, production switched on :80)
 - [x] Rollback test — green → blue between two one-second samples, zero failed requests
-


2. Part 1 — Auto Scaling for a 3-Tier Application

Student: Joyanta Sarker Joy · **Course:** Mastering DevOps, Batch 13 · **Region:** Asia Pacific (Mumbai) `ap-south-1`

1. Objective

Take a backend that runs on **one** EC2 instance — one server, one point of failure — and turn it into a **highly available, auto-scaling, self-healing** 3-tier architecture. When a server dies, a new one replaces it automatically. When traffic rises, more servers appear. When it falls, they go away. No human has to touch anything.

2. Architecture



The three tiers

Tier	What it is	Built with
Tier 1 — Presentation / Edge	The public front door. The only thing the internet can reach.	Application Load Balancer <code>bitopi-alb</code>
Tier 2 — Application (Backend)	The servers that run the code. This is the tier that auto-scales.	Auto Scaling Group <code>bitopi-app-asg</code> of EC2 instances running a Node.js app on port 3000
Tier 3 — Data	The database every backend server reads from and writes to.	MySQL (MariaDB) on EC2 <code>bitopi-db-server</code>

How a request travels

1. **Browser → ALB** on port **80**. The ALB's DNS name is the one public address for the whole system.
2. **ALB → a healthy backend EC2** on port **3000**. The ALB only chooses servers that pass its health check (`GET /health → 200`).
3. **Backend EC2 → MySQL** on port **3306**, using the database server's private IP `10.0.13.216`.
4. The answer travels back along the same path.

Why two Availability Zones

An Availability Zone is a physically separate data centre. Every layer here exists in **both** `ap-south-1a` and `ap-south-1b`: the ALB has a node in each, the Auto Scaling Group spreads instances across both, and the subnets are paired. If one data centre has a problem, the other keeps serving. That is what "highly available" means in practice.

3. Design decisions and account constraints

This was built in a **restricted course account**. Three limits shaped the design, and the write-up is honest about each one.

Constraint	What it blocked	What we did instead
Cannot create RDS resources (<code>rds:CreateDBSubnetGroup</code> denied)	A managed database	Installed MySQL (MariaDB 10.5) on a <code>t3.micro</code> EC2, fully auto-configured by a first-boot script. Same port, same firewall rule, same app code — the app cannot tell the difference.
No IAM role creation, no SSM	Session Manager, Parameter Store, instance profiles	SSH with a key pair for access; the app reads its config from environment variables written at boot.
No CloudShell	AWS CLI in the browser	Everything was done in the console; permissions were confirmed by building.

Two deliberate cost/complexity choices:

- **No NAT Gateway.** A NAT Gateway costs ~US\$32/month. Instead the backend and database servers sit in the **public** subnets with a public IP (so they can download packages), while the **security groups** guarantee that nothing on the internet can reach the app port or the database port. The two private subnets exist and are the correct home for the data tier in production (with RDS, or with a NAT Gateway / VPC endpoints for outbound access).
- **MySQL on one EC2, not Multi-AZ.** Adequate for the lab; in production this would be RDS Multi-AZ with a standby in the second AZ, as drawn on the diagram.

4. Build steps

Each step lists **what** was created, **why** it exists, **when you would use it**, the exact settings, and the evidence screenshot.


Step 1 — VPC and subnets

What: a private network `bitopi-vpc` (`10.0.0.0/16`) with two public and two private subnets across two AZs, an Internet Gateway, and route tables.


Why: everything else lives inside this network. Public subnets have a route to the Internet Gateway; private subnets do not — that is the only difference, and it is what makes a subnet "public" or "private".

When you'd use it: every AWS workload that is not a throw-away test starts with its own VPC. Never build production on the default VPC.

Setting	Value
Method	VPC console → <i>Create VPC</i> → " VPC and more "
Name prefix	<code>bitopi</code>
IPv4 CIDR	<code>10.0.0.0/16</code>
Availability Zones	2 — <code>ap-south-1a</code> , <code>ap-south-1b</code>
Public subnets	<code>bitopi-subnet-public1-ap-south-1a</code> <code>10.0.0.0/20</code> · <code>bitopi-subnet-public2-ap-south-1b</code> <code>10.0.16.0/20</code>
Private subnets	<code>bitopi-subnet-private1-ap-south-1a</code> <code>10.0.128.0/20</code> · <code>bitopi-subnet-private2-ap-south-1b</code> <code>10.0.144.0/20</code>
NAT gateways	None (cost)
DNS hostnames / resolution	Enabled
Auto-assign public IPv4	Enabled on both public subnets (so Auto Scaling instances get a public IP for package downloads)



Resource map: one VPC, four subnets, one public route table shared by both public subnets (with a route to `bitopi-igw`), one private route table per private subnet (no internet route).



Step 2 — Security groups (the firewalls)

What: three security groups, one per tier. Each one only trusts the tier directly in front of it.

Why: this is *defence in depth*. Even though the backend and database servers have public IPs, the firewall rules mean the only thing that can reach port 3000 is the load balancer, and the only thing that can reach port 3306 is a backend server. The rules reference **other security groups**, not IP addresses — so they keep working no matter how many instances the Auto Scaling Group launches or replaces.

When you'd use it: always. Chaining security groups by reference is the standard AWS pattern for multi-tier apps.

Security group	ID	Inbound rules
bitopi-alb-sg	sg-04c8272fbbeb0e769	HTTP 80 from 0.0.0.0/0 · HTTPS 443 from 0.0.0.0/0
bitopi-backend-sg	sg-0029472c273fdd274	TCP 3000 from bitopi-alb-sg · SSH 22 from admin
bitopi-db-sg	sg-01883e3876cbf420d	MySQL 3306 from bitopi-backend-sg · SSH 22 from admin

Outbound on all three: default (all traffic). Security groups are **stateful**: when the backend opens a connection *into* the database, the reply is automatically allowed back out, so only inbound rules need to be written.

A note on SSH. The SSH rules were left at 0.0.0.0/0 because the admin's public IP changes between networks and a "My IP" rule locked him out. Access still requires the bitopi-key.pem private key, so password-guessing is impossible. In production, SSH would be restricted to a bastion host or replaced with SSM Session Manager (not available in this account).



Security Groups (1/7)

Name	Security group ID	Security group name	VPC ID	Description	Owner	Inbound rules count
-	sg-0adfd480cc1b041cd4	default	vpc-073640928ee1bc341	default VPC security group	448513989308	1 Permission entry
-	sg-0fb9e957cb102e0ca	default	vpc-086f2b37fc9b2a2c	default VPC security group	448513989308	1 Permission entry
-	sg-03d99acafb09c2146	default	vpc-0ab019c806b8d4938	default VPC security group	448513989308	1 Permission entry
-	sg-0ff49c85295d39d4e	default	vpc-0e50de8c43d883c3a	default VPC security group	448513989308	2 Permission entries
1_Bitopi-ALB-SG	sg-04c8272fbbeb0e769	bitopi-alb-sg	vpc-0ab019c806b8d4938	Allow HTTP/HTTPS from internet to ALB	448513989308	2 Permission entries
2_Bitopi-Backend-SG	sg-0029472c273fdd274	bitopi-backend-sg	vpc-0ab019c806b8d4938	Allow app port from ALB and SSH from ...	448513989308	2 Permission entries
3_Bitopi-DB	sg-01883e3876cbf420d	bitopi-db-sg	vpc-0ab019c806b8d4938	Allow MySQL from backend only	448513989308	2 Permission entries

sg-01883e3876cbf420d - bitopi-db-sg

Details

Security group name bitopi-db-sg	Security group ID sg-01883e3876cbf420d	Description Allow MySQL from backend only	VPC ID vpc-0ab019c806b8d4938
Owner 448513989308	Inbound rules count 2 Permission entries	Outbound rules count 1 Permission entry	



Step 3 — Database tier (MySQL on EC2)

What: one `t3.micro` running Amazon Linux 2023. Its first-boot **user-data** script installs MariaDB, makes it listen on the network, creates the `appdb` database and an `appuser` account, and seeds two tables. Nobody ever logged in to set it up.


Why: the app needs a shared place to keep data. Every backend server — however many the Auto Scaling Group creates — connects to this one database, so all of them see the same data.

When you'd use it: self-hosting a database on EC2 is what you do when a managed service is unavailable (as here) or when you need something RDS doesn't offer. Otherwise, use RDS.

Setting	Value
Name / ID	bitopi-db-server / i-0496778ef511974dd
AMI · type	Amazon Linux 2023 · t3.micro
Subnet	bitopi-subnet-public1-ap-south-1a (public IP enabled for package install)
Security group	bitopi-db-sg
Key pair	bitopi-key
Private IP	10.0.13.216 — this is what the app connects to
Public IP	13.233.34.184 (admin SSH only)
User-data	scripts/db-server-userdata.sh

The script, in plain words:

```
dnf install -y mariadb105-server          # install MySQL-compatible server
systemctl enable --now mariadb            # start it now and on every boot
# make it listen on all interfaces (not just localhost)
printf '[mysqld]\nbind-address=0.0.0.0\n' > /etc/my.cnf.d/99-app.cnf
systemctl restart mariadb
mysql <<'SQL'                             # create database, user, tables, seed row
CREATE DATABASE IF NOT EXISTS appdb;
CREATE USER IF NOT EXISTS 'appuser'@'%' IDENTIFIED BY '...';
GRANT ALL PRIVILEGES ON appdb.* TO 'appuser'@'%';
...
SQL
```



Step 4 — Launch Template (the recipe for every app server)

What: `bitopi-app-lt` — a saved blueprint that says exactly how to build one backend server: which OS image, which size, which key, which firewall, and a first-boot script that installs and starts the app.

Why: the Auto Scaling Group must be able to create servers **without a human**. It can only do that if it has a recipe. Every instance stamped from this template is identical.

When you'd use it: any time more than one copy of a server will ever exist — Auto Scaling Groups, spot fleets, or simply "launch another one exactly like the last one".

Setting	Value
Name / ID	<code>bitopi-app-lt</code> / <code>lt-072cb16544d51e361</code> (v1)
AMI · type	Amazon Linux 2023 (<code>ami-066c4849e6b3a1e3d</code>) · <code>t3.micro</code>
Key pair	<code>bitopi-key</code>
Security group	<code>bitopi-backend-sg</code>
Subnet	not included — the ASG decides, so it can spread across AZs
User-data	<code>scripts/app-server-userdata.sh</code>


What the user-data does on every new server, in order:

1. `dnf install -y nodejs` — install Node.js.
2. Ask the **instance metadata service** for its own instance-ID and AZ (IMDSv2: get a token, then query `169.254.169.254`).
3. Write `/opt/app/app.js` and `package.json` (`app/`).
4. Write `/opt/app/.env` with `PORT=3000` , its instance-ID and AZ, and the database connection (`DB_HOST=10.0.13.216` , `appuser` , `appdb` , `3306`).
5. `npm install` — fetch `express` and `mysql2` .

6. Register a **systemd service** `bitopi-app` so the app starts on boot and restarts if it crashes.

The app itself (`app/app.js`) serves two routes:

- `GET /health` → `200 {"status":"healthy", ...}`. **Deliberately shallow** — it only proves the process is alive. It does *not* touch the database, so a brief DB blip cannot make the load balancer mark the whole fleet unhealthy.
- `GET /` → an HTML page showing which instance answered, its AZ, whether the database is reachable, the seed row read from MySQL, and a visit counter that *writes* a row to MySQL on every load. One page proves all three tiers are talking.



Step 5 — Target Group and Application Load Balancer

What: the Target Group `bitopi-app-tg` is the ALB's *list of servers* plus the *health-check rule*. The ALB `bitopi-alb` is the public entry point that listens on port 80 and forwards to that group.

Why: users need **one** stable address while the servers behind it come and go. The ALB provides it, spreads traffic across every healthy server, and quietly stops sending traffic to any server that fails its health check.

When you'd use it: any web application with more than one server, or any application that must survive a server failure without users noticing.

Target group <code>bitopi-app-tg</code>	
Target type · protocol : port	Instances · HTTP : 3000
VPC	<code>bitopi-vpc</code>
Health check	HTTP GET /health , success code <code>200</code>
Interval · timeout	15 s · 5 s
Healthy · unhealthy threshold	2 · 2 (so a server flips state after 30 s of consistent results)
Registered targets at creation	none — the ASG registers them

Load balancer bitopi-alb	
Type · scheme	Application · Internet-facing
Subnets	both public subnets (one node per AZ)
Security group	bitopi-alb-sg
Listener	HTTP : 80 → forward to bitopi-app-tg
DNS name	bitopi-alb-1957667478.ap-south-1.elb.amazonaws.com

The screenshot shows the AWS Management Console interface for the 'bitopi-app-tg' target group. A green notification bar at the top states: 'Successfully created the target group: bitopi-app-tg. Anomaly detection is automatically applied to all registered targets. Results can be viewed in the Targets tab.' The left sidebar shows the navigation menu with 'EC2' selected. The main content area displays the 'Details' tab for the target group, showing its ID, target type (Instance), IP address type (IPv4), protocol (HTTP), and port (3000). Below the details, there are status indicators for 'Total targets' (0), 'Healthy' (0), 'Unhealthy' (0), 'Unused' (0), 'Initial' (0), and 'Draining' (0). The 'Targets' tab is active, showing a table with columns for Instance ID, Name, Port, Zone, Health status, Health status details, Administrative o..., Override details, and Launch time. A message indicates 'No registered targets' and provides a 'Register targets' button.

The screenshot shows the AWS Management Console interface for the 'bitopi-alb' load balancer. A green notification bar at the top states: 'Successfully created load balancer: bitopi-alb. It might take a few minutes for your load balancer to fully set up and route traffic. Targets will also take a few minutes to complete the registration process and pass initial health checks.' Below this, a blue notification bar mentions 'Introducing new Capacity Utilization metric for ALB'. The left sidebar shows the navigation menu with 'EC2' selected. The main content area displays the 'Details' tab for the load balancer, showing its type (Application), status (Provisioning), VPC, hosted zone, availability zones, and load balancer IP address type. The 'Listeners and rules' tab is also visible, showing a table with columns for Protocol/Port, Default action, Rules, ARN, Security policy, Default SSL/TLS certificate, mTLS, and Trust store. A message indicates 'No registered targets' and provides a 'Register targets' button.

The screenshot shows the 'Listeners and rules' configuration for the load balancer. The 'Listeners and rules (1)' tab is active, showing a table with columns for Protocol/Port, Default action, Rules, ARN, Security policy, Default SSL/TLS certificate, mTLS, and Trust store. A single listener is listed with the protocol 'HTTP:80' and the default action 'Forward to target group bitopi-app-tg (100%)'. The 'Rules' column shows '1 rule' and the 'ARN' column shows 'ARN'. The 'Security policy' column shows 'Not applicable', the 'Default SSL/TLS certificate' column shows 'Not applicable', the 'mTLS' column shows 'Not applicable', and the 'Trust store' column shows 'Not applicable'.

Step 6 — Auto Scaling Group

What: bitopi-app-asg — the manager. It keeps the number of running app servers between **min 2** and **max 4**, launches them from the Launch Template into both public subnets, registers each one with the Target Group, and replaces any that become unhealthy.



Why: this single component delivers all three assignment goals. **High availability** — two servers in two AZs. **Self-healing** — it watches health and replaces failures. **Auto scaling** — with a policy attached (Step 8) it adds and removes servers on demand.

When you'd use it: whenever a stateless service needs to stay up and cope with variable load. Web/API backends are the classic case.

Setting	Value
Launch template	bitopi-app-lt , version <i>Latest</i>
Subnets	bitopi-subnet-public1-ap-south-1a + bitopi-subnet-public2-ap-south-1b
Load balancing	attached to target group bitopi-app-tg
Health checks	EC2 and Elastic Load Balancing (so a failed <code>/health</code> counts as unhealthy, not just a stopped VM)
Health check grace period	300 s — a new server gets 5 minutes to install Node and start before its health is judged
Capacity	min 2 · desired 2 · max 4
Tag	Name = bitopi-app-server (propagated to every instance)

Result: two instances launched, one per AZ, and both passed the health check within ~4 minutes. Opening the ALB's DNS name in a browser served the app page with **Database: CONNECTED** — browser → ALB → EC2 → MySQL, all working. Refreshing the page shows different instance IDs as the ALB rotates between servers.

The screenshot displays the AWS Management Console interface for an Auto Scaling group. The left sidebar shows the navigation menu with categories like EC2, Instances, Images, Elastic Block Store, Network & Security, Load Balancing, and Auto Scaling. The main content area is titled 'Auto Scaling groups (1/1) info' and includes a table listing the group 'bitopi-app-asg'. Below this, the 'Details' tab is selected, showing the 'bitopi-app-asg Capacity overview' and 'Launch template' details. The capacity overview shows a desired capacity of 2, scaling limits of 2-4, and a status of 'At desired capacity'. The launch template details show the launch template 'lt-072cb16544d51e361', AMI ID 'ami-066c4b49e6b3a1e3d', instance type 't3.micro', and security group 'sg-0029472c273fdd274'. The console footer includes the AWS logo, navigation links, and copyright information for 2026.



Step 7 — Self-healing test

Test: with two healthy instances running, one (`i-0022dbdb97f0a99e3`, in `ap-south-1a`) was **terminated by hand** from the EC2 console — simulating a hardware failure or a crashed server.

What happened, from the ASG activity log:

Time (UTC+6)	Event
11:03:08	"Terminating EC2 instance i-0022dbdb97f0a99e3 ... taken out of service in response to an EC2 health check ... Waiting For ELB Connection Draining"
11:03:10	"Launching a new EC2 instance i-0b762fab68c3abe53 ... in response to an unhealthy instance needing to be replaced"
~11:06	Target group back to 2 healthy — the replacement placed in <code>ap-south-1a</code> , restoring the AZ balance

Two seconds from detection to replacement. During the rebuild the app URL kept working, served by the surviving instance in `ap-south-1b` — **zero downtime**. The replacement has a **different instance ID**: it is a brand-new server built from the Launch Template, not a restart.

Connection draining means the ALB let in-flight requests to the dying server finish before cutting it off.


The screenshot shows the AWS Management Console for an Auto Scaling Group named 'bitopi-app-asg'. The 'Activity' tab is selected, showing a list of events. The events include:

- Successful**: Launching a new EC2 instance i-0b762fab68c3abe53. At 2026-09-25T05:03:08Z an instance was launched in response to an unhealthy instance needing to be replaced.
- Connection draining in progress**: Terminating EC2 instance i-0022dbdb97f0a99e3 - Waiting For ELB Connection Draining. At 2026-09-25T05:03:08Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.
- Successful**: Launching a new EC2 instance i-07252287a3f116ae2. At 2026-09-25T04:38:46Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 2. At 2026-09-25T04:38:47Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 2.
- Successful**: Launching a new EC2 instance i-0022dbdb97f0a99e3. At 2026-09-25T04:38:46Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 2. At 2026-09-25T04:38:47Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 2.

The screenshot shows the Bitopi 3-Tier Application web interface. The application is titled 'Bitopi 3-Tier Application' and is running on 'BACKEND TIER - AUTO SCALING'. The interface displays the following details:

- Served by instance**: i-07252287a3f116ae2
- Availability Zone**: ap-south-1b
- Hostname**: ip-10-0-23-139.ap-south-1.compute.internal
- Database (Tier 3)**: CONNECTED
- DB message**: Bitopi 3-Tier App - data served live from MySQL on EC2
- Total visits recorded**: 5

A note at the bottom states: 'Refresh — the load balancer may route you to a different instance. Every visit is written to MySQL, proving the Backend — Database link.'



Step 8 — Scaling policy and load test

Policy: `bitopi-cpu-target-tracking` — **Target tracking** on **Average CPU utilization = 30%**, instance warm-up 60 s, scale-in enabled.

Why target tracking: you give it one number and AWS derives everything else — it creates two CloudWatch alarms (`AlarmHigh` , `AlarmLow`) and works out how many instances keep the average near the target. It is the *thermostat* model. The alternative, **step scaling**, makes you write the rules yourself ("if CPU > 60% add 1; if > 80% add 2"). Target tracking is the recommended default for web backends.

How the maths works (this predicted the result exactly):

$$\text{desired} = \text{ceil}(\text{current instances} \times \text{current metric} \div \text{target}) = \text{ceil}(2 \times 50\% \div 30\%) = \text{ceil}(3.33) = 4$$

Load generation: SSH into a backend instance and pin every CPU core with a dependency-free loop that stops itself after 15 minutes:

```
for i in $(seq 1 $(nproc)); do timeout 900 bash -c 'while :; do :; done' & done
top -bn1 | head -5          # showed: 86.1 us, 13.9 sy, 0.0 id → 100% busy
```

Stop early with `kill -f "timeout 900"`.

Scale-out, from the activity log:

Time	Event
~11:27	Load started on <code>i-07252287a3f116ae2</code> (<code>ip-10-0-23-139</code>) — 0.0% idle
11:34:42	" <i>monitor alarm TargetTracking-...-AlarmHigh in state ALARM triggered policy bitopi-cpu-target-tracking changing the desired capacity from 2 to 4</i> " — launched <code>i-0f33782fbae77fa57</code> and <code>i-0791db15f912ef398</code>
11:35:47	4 instances InService, all Healthy , two per AZ

Scale-in (load stopped ~11:45; target tracking waits ~15 minutes of sustained low CPU before removing anything, to prevent "flapping"):

Time	Event
12:01:20	" <i>AlarmLow in state ALARM ... changing the desired capacity from 4 to 3</i> " — <code>i-07252287a3f116ae2</code> selected, connection draining
12:01:40	" <i>... from 3 to 2</i> " — <code>i-0b762fab68c3abe53</code> selected, connection draining
12:04	Desired 2; one instance left in each AZ (the ASG removed one from each zone, keeping the balance)

The screenshot shows the AWS Management Console for the 'bitopi-app-asg' Auto Scaling Group. The 'Automatic scaling' tab is active, showing the following configuration:

- Desired capacity:** 2
- Scaling limits:** 2 - 4
- Desired capacity type:** Units (number of instances)
- Status:** At desired capacity
- Date created:** Fri Sep 25 2026 10:38:46 GMT+0600 (Bangladesh Standard Time)

The 'Automatic scaling' section includes a description of scaling policies and a list of dynamic scaling policies. The 'bitopi-cpu-target-tracking' policy is highlighted, showing its configuration:

- Policy type:** Target tracking scaling
- Enabled or disabled:** Enabled
- Execute policy when:** As required to maintain Average CPU utilization at 30
- Take the action:** Add or remove capacity units as required
- Instances need:** 60 seconds to warm up before including in metric
- Scale in:** Enabled

The right sidebar provides additional information about dynamic scaling policies, including a note about capacity units and their relationship to instance types.

bitopi-app-asg

bitopi-app-asg Capacity overview

arn:aws:autoscaling:ap-south-1:448513989308:autoScalingGroup:5fb9be00-6794-4c7f-8fb6-cab96bb2f0d:autoScalingGroupName/bitopi-app-asg

Desired capacity: 4 Scaling limits: 2 - 4 Desired capacity type: Units (number of instances) Status: At desired capacity

Date created: Fri Sep 25 2026 10:38:46 GMT+0600 (Bangladesh Standard Time)

Details | Integrations | Automatic scaling | **Instance management** | Instance refresh | Activity | Monitoring | Tags

Instances (4)

Instance ID	Lifecycle	Instance type	Weighted capac...	Launch templ...	Availability Zone	Health status	Protected from
i-0725287a3f116ae2	InService	t3.micro	-	bitopi-app-it	ap-s1-a23 (ap-sout...	Healthy	
i-0791db15f912ef398	InService	t3.micro	-	bitopi-app-it	ap-s1-a23 (ap-sout...	Healthy	
i-0b762fab68c3abe53	InService	t3.micro	-	bitopi-app-it	ap-s1-a21 (ap-sout...	Healthy	
i-0f33782fbac77fa57	InService	t3.micro	-	bitopi-app-it	ap-s1-a21 (ap-sout...	Healthy	

Instance lifecycle policy for lifecycle hooks

Controls instance behavior when an instance transitions through its lifecycle states.

Termination hook abandon behavior: Terminate (default)

Lifecycle hooks (0)

Dynamic scaling policies

Amazon EC2 Auto Scaling can add more instances (referred to as *scaling out*) to deal with high demand at peak times, and run fewer instances (referred to as *scaling in*) to reduce costs during periods of low utilization.

When you create a target tracking scaling policy, Amazon EC2 Auto Scaling automatically increases and decreases capacity in response to varying usage levels. For example, a target tracking scaling policy might have a target CPU value of 50 percent. Amazon EC2 Auto Scaling then launches and terminates EC2 instances as required to keep the aggregated CPU usage across all instances in your group at 50 percent.

With step scaling and simple scaling, you must create alarms in Amazon CloudWatch, and then define two policies, one for scaling out and the other for scaling in. Step scaling can make bigger or smaller size adjustments based on the metric value, while simple scaling always makes the same size adjustment.

Note

Capacity units takes into account the unit of measurement that you chose when configuring the Auto Scaling group. By default, one capacity unit represents one instance. If you chose to use multiple instance types and

bitopi-app-asg

Details | Integrations | Automatic scaling | Instance management | Instance refresh | **Activity** | Monitoring | Tags

Activity notifications (0)

No notifications are currently specified

Activity history (8)

Status	Description	Cause	Start time	End time
Connection draining in progress	Terminating EC2 instance i-0b762fab68c3abe53 - Waiting For ELB Connection Draining.	At 2026-09-25T06:01:40Z a monitor alarm TargetTracking-bitopi-app-asg-AlarmLow-50Sc4dd4-18cf-49cb-870f-be98ca928766 in state ALARM triggered policy bitopi-cpu-target-tracking changing the desired capacity from 3 to 2. At 2026-09-25T06:01:46Z an instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 3 to 2. At 2026-09-25T06:01:46Z instance i-0b762fab68c3abe53 was selected for termination.	2026 September 25, 12:01:46 PM +06:00	
Connection draining in progress	Terminating EC2 instance i-0725287a3f116ae2 - Waiting For ELB Connection Draining.	At 2026-09-25T06:01:20Z a monitor alarm TargetTracking-bitopi-app-asg-AlarmLow-50Sc4dd4-18cf-49cb-870f-be98ca928766 in state ALARM triggered policy bitopi-cpu-target-tracking changing the desired capacity from 4 to 3. At 2026-09-25T06:01:24Z an instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 4 to 3. At 2026-09-25T06:01:25Z instance i-0725287a3f116ae2 was selected for termination.	2026 September 25, 12:01:25 PM +06:00	
Successful	Launching a new EC2 instance i-0f33782fbac77fa57	At 2026-09-25T05:34:26Z a monitor alarm TargetTracking-bitopi-app-asg-AlarmHigh-63055eda-2f84-44a3-a08a-404c4da0190c in state ALARM triggered policy bitopi-cpu-target-tracking changing the desired capacity from 2 to 4. At 2026-09-25T05:34:40Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 2 to 4.	2026 September 25, 11:34:42 AM +06:00	2026 September 25, 11:35:47 AM +06:00
Successful	Launching a new EC2 instance i-0791db15f912ef398	At 2026-09-25T05:34:26Z a monitor alarm TargetTracking-bitopi-app-asg-AlarmHigh-63055eda-2f84-44a3-a08a-404c4da0190c in state ALARM triggered policy bitopi-cpu-target-tracking changing the desired capacity from 2 to 4. At 2026-09-25T05:34:40Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 2 to 4.	2026 September 25, 11:34:42 AM +06:00	2026 September 25, 11:35:46 AM +06:00

bitopi-app-asg

bitopi-app-asg Capacity overview

arn:aws:autoscaling:ap-south-1:448513989308:autoScalingGroup:5fb9be00-6794-4c7f-8fb6-cab96bb2f0d:autoScalingGroupName/bitopi-app-asg

Desired capacity: 2 Scaling limits: 2 - 4 Desired capacity type: Units (number of instances) Status: Updating capacity

Date created: Fri Sep 25 2026 10:38:46 GMT+0600 (Bangladesh Standard Time)

Details | Integrations | Automatic scaling | **Instance management** | Instance refresh | Activity | Monitoring | Tags

Instances (4)

Instance ID	Lifecycle	Instance type	Weighted capacity	Launch template...	Availability Zone	Health status	Protected from
i-0725287a3f116ae2	Terminating	t3.micro	-	bitopi-app-it	ap-s1-a23 (ap-south...	Healthy	
i-0791db15f912ef398	InService	t3.micro	-	bitopi-app-it	ap-s1-a23 (ap-south...	Healthy	
i-0b762fab68c3abe53	Terminating	t3.micro	-	bitopi-app-it	ap-s1-a21 (ap-south...	Healthy	
i-0f33782fbac77fa57	InService	t3.micro	-	bitopi-app-it	ap-s1-a21 (ap-south...	Healthy	

Instance lifecycle policy for lifecycle hooks

Controls instance behavior when an instance transitions through its lifecycle states.

Termination hook abandon behavior: Terminate (default)

Lifecycle hooks (0)

5. Health checks — the two layers explained

There are two different health checks and they do different jobs.

	ALB / Target Group health check	Auto Scaling Group health check
Question it asks	"Should I send this server traffic <i>right now?</i> "	"Should this server be <i>replaced?</i> "
How	GET <code>/health</code> every 15 s; 2 consecutive failures → <i>unhealthy</i>	EC2 status checks plus the ELB result (we turned this on)
Consequence	Server is taken out of rotation; traffic goes to the others	Server is terminated and a new one launched from the template
Safety valve	—	300 s grace period so a booting server isn't killed for being slow

Turning on *ELB health checks* in the ASG is what links the two: an app that hangs (process alive, but `/health` not answering) fails the ALB check, the ASG hears about it, and the server is replaced. Without that setting the ASG would only replace servers whose VM died.

6. Submission checklist → evidence

Required item	Where
Architecture diagram	<code>diagrams/part1_architecture.png</code>
Launch Template configuration	P1-07
ALB & Target Group configuration	P1-09, P1-10, P1-11
Auto Scaling Group configuration	P1-12
Health check configuration	P1-09 (target group), P1-12 (ASG)
Security groups	P1-02, P1-03, P1-04
Self-healing test evidence	P1-15, P1-16, P1-17
Scaling policy configuration	P1-18
Scale-out / scale-in evidence	P1-19 ... P1-23
Application running through the ALB	P1-14
Short documentation	this file

7. Resource inventory

Resource	Name	ID / address
VPC	bitopi-vpc	vpc-0ab019c806b84a938 · 10.0.0.0/16
Internet gateway	bitopi-igw	—
Public subnets	...public1-ap-south-1a / ...public2-ap-south-1b	subnet-053a27c29e6c271ad / subnet-0d4371ebe7ff2dcf8
Security groups	bitopi-alb-sg / bitopi-backend-sg / bitopi-db-sg	sg-04c8272fbee0e769 / sg-0029472c273fdd274 / sg-01883e3876cbf420d
Key pair	bitopi-key	bitopi-key.pem (never committed)
Database server	bitopi-db-server	i-0496778ef511974dd · private 10.0.13.216
Launch template	bitopi-app-lt	lt-072cb16544d51e361
Target group	bitopi-app-tg	HTTP : 3000, /health
Load balancer	bitopi-alb	bitopi-alb-1957667478.ap-south-1.elb.amazonaws.com
Auto Scaling group	bitopi-app-asg	min 2 / desired 2 / max 4
Scaling policy	bitopi-cpu-target-tracking	avg CPU 30 %

8. Cleanup

The infrastructure is left running because **Part 2 (CI/CD) deploys onto this same Auto Scaling Group**. The cleanup checklist is in the final section of the Part 2 document.


3. How the Servers Talk to Each Other

This guide follows one request through the whole system and shows, at every hop, **who** talks to **whom**, on **which port**, **which firewall rule** allows it, and **which command** you can run to see it happen. Every command below was run on this project's real servers.

Two ideas make all of it work:

- **Security groups reference each other, not IP addresses.** `bitopi-backend-sg` allows port 3000 from `bitopi-alb-sg` — meaning "from anything wearing the ALB's security group". That rule never has to change when instances come and go.
- **Security groups are stateful.** If a connection is allowed *in*, its reply is automatically allowed *out*. So we only write inbound rules, and every "vice versa" below is the reply traffic of an allowed inbound connection — no extra rule needed.

The map



The backend servers **never talk to each other**. They are stateless twins; the only thing they share is the database. That is deliberate — it is what lets the Auto Scaling Group add or remove them freely.

Hop 1 — Browser → Load Balancer (port 80)

From → to	your laptop → bitopi-alb
Port / protocol	80 / HTTP
Address	bitopi-alb-1957667478.ap-south-1.elb.amazonaws.com (a DNS name; AWS maps it to one ALB node per AZ)
Allowed by	bitopi-alb-sg inbound: HTTP 80 from 0.0.0.0/0

The ALB is the **only** component with a rule that says "from anywhere". Everything behind it is invisible to the internet.

See it (from your laptop, PowerShell):

```
# Which IPs does the ALB name resolve to? (one per AZ – that's the multi-AZ front door)
nslookup bitopi-alb-1957667478.ap-south-1.elb.amazonaws.com

# Send a request and see only the response headers
curl.exe -I http://bitopi-alb-1957667478.ap-south-1.elb.amazonaws.com/

# Hit it 6 times; the instance ID in the page changes as the ALB rotates between servers
1..6 | ForEach-Object { (curl.exe -s http://bitopi-alb-1957667478.ap-south-1.elb.amazonaws.com/ |
Select-String 'i-0[0-9a-f]+').Matches.Value }

# Prove the backend is NOT reachable directly on port 3000 from the internet (times out - the
firewall works)
curl.exe -m 5 http://<any-backend-public-ip>:3000/
```

Hop 2 — Load Balancer → Backend EC2 (port 3000) — and back

From → to	ALB node → a backend instance
Port / protocol	3000 / HTTP
Address	the instance's private IP (e.g. 10.0.23.139), registered in target group bitopi-app-tg
Allowed by	bitopi-backend-sg inbound: TCP 3000 from bitopi-alb-sg
Vice versa	the reply is the same TCP connection → allowed by statefulness

This hop carries two kinds of traffic:

1. **User requests** — forwarded to a healthy instance chosen round-robin.
2. **Health checks** — every 15 s the ALB itself calls GET /health on every registered instance. Two failures in a row → the instance is marked *unhealthy* and stops receiving user traffic. Two passes → healthy again.

See it (SSH into any backend instance):

```
# 1) Is the app listening on 3000?
sudo ss -tlnp | grep 3000
#   LISTEN 0 511 *:3000 *:~ users:(("node",pid=...))

# 2) Call the two routes exactly as the ALB does
curl -s http://localhost:3000/health # →
{"status":"healthy","instance":"i-...", "az":"ap-south-1b"}
curl -s http://localhost:3000/ | head -20 # the HTML page

# 3) Watch the ALB's health checks arriving in real time (the source IPs are the ALB's private IPs)
sudo journalctl -u bitopi-app -f
# or, at the network level:
sudo tcpdump -ni any port 3000 -c 10

# 4) Which security group is on this instance, and what does it allow?
TOKEN=$(curl -s -X PUT http://169.254.169.254/latest/api/token -H "X-aws-ec2-metadata-token-ttl-seconds: 60")
curl -s -H "X-aws-ec2-metadata-token: $TOKEN" http://169.254.169.254/latest/meta-data/security-groups
#   bitopi-backend-sg
```

Why the health check is shallow. `/health` only proves the Node process is alive; it does not query the database. If it did, a 10-second database hiccup would make every instance fail at once, the ALB would have nowhere to send traffic, and the ASG would start replacing perfectly good servers. Keep liveness checks cheap; show deeper status on a separate page (which is what `/` does).

Hop 3 — Backend EC2 → Database (port 3306) — and back

From → to	any backend instance → <code>bitopi-db-server</code>
Port / protocol	3306 / MySQL protocol
Address	private IP 10.0.13.216 — never the public one; traffic stays inside the VPC
Allowed by	<code>bitopi-db-sg</code> inbound: MySQL 3306 from <code>bitopi-backend-sg</code>
Vice versa	query results come back on the same connection → statefulness. The database never initiates a connection to anything.

In the app this hop is one line of configuration (`DB_HOST=10.0.13.216` in `/opt/app/.env`) and a few lines of code using the `mysql2` driver:

```
const conn = await mysql.createConnection({ host: process.env.DB_HOST, user: 'appuser', ... });
await conn.execute('INSERT INTO visits (hostname) VALUES (?)', [os.hostname()]);
const [rows] = await conn.query('SELECT info FROM app_info WHERE id=1');
```

See it (SSH into a backend instance):

```
# 1) Can I even reach the database port? (network + firewall test, no credentials needed)
nc -zv 10.0.13.216 3306
# Connection to 10.0.13.216 3306 port [tcp/mysql] succeeded!

# 2) Install the MySQL client and run a real query as the app user
sudo dnf install -y mariadb105
mysql -h 10.0.13.216 -u appuser -p'Bitopi#Db2026' appdb -e "SELECT * FROM app_info; SELECT COUNT(*) AS visits FROM visits;"

# 3) See the app's own connection config
cat /opt/app/.env

# 4) Watch the live connections from this server to the DB
sudo ss -tn | grep 10.0.13.216:3306
```

See it from the other side (SSH into `bitopi-db-server`):

```
# The database listens on all interfaces (we set bind-address=0.0.0.0)
sudo ss -tlnp | grep 3306

# Who is connected right now? (the Host column shows backend private IPs)
sudo mysql -e "SELECT user, host, db, command FROM information_schema.processlist;"

# Every row the backends have written
sudo mysql -e "SELECT * FROM appdb.visits ORDER BY id DESC LIMIT 10;"
```

Proof the firewall is doing its job: run `nc -zv 10.0.13.216 3306` from the **ALB's** point of view is impossible (ALBs can't run commands), but from your **laptop** it times out — port 3306 is not open to the internet, only to `bitopi-backend-sg`.

Hop 4 — Backend EC2 → Instance Metadata Service (IMDSv2)

From → to	an instance → <code>169.254.169.254</code> (a link-local address that exists inside every EC2)
Port / protocol	80 / HTTP, IMDSv2 = token first, then queries
Allowed by	nothing to allow — it never leaves the host; security groups don't apply

This is how a freshly launched server learns *who it is* without anyone telling it. The user-data script uses it to stamp the instance-ID and AZ into the app's config — which is why the web page can show "Served by instance `i-...` in `ap-south-1b`".

```
# Step 1: get a short-lived token (IMDSv2 requires this; the instance was launched with "IMDSv2:
Required")
TOKEN=$(curl -s -X PUT "http://169.254.169.254/latest/api/token" -H "X-aws-ec2-metadata-token-
ttl-seconds: 300")

# Step 2: ask questions with the token
curl -s -H "X-aws-ec2-metadata-token: $TOKEN" http://169.254.169.254/latest/meta-data/instance-id
curl -s -H "X-aws-ec2-metadata-token: $TOKEN" http://169.254.169.254/latest/meta-
data/placement/availability-zone
curl -s -H "X-aws-ec2-metadata-token: $TOKEN" http://169.254.169.254/latest/meta-data/local-ipv4
curl -s -H "X-aws-ec2-metadata-token: $TOKEN" http://169.254.169.254/latest/meta-data/public-ipv4
curl -s -H "X-aws-ec2-metadata-token: $TOKEN" http://169.254.169.254/latest/user-data | head -20
# the script that built this server
```

Hop 5 — Backend EC2 → the Internet (package downloads)

From → to	an instance → <code>dnf mirrors</code> , <code>registry.npmjs.org</code>
Path	instance's public IP → public route table → <code>bitopi-igw</code> → internet
Allowed by	outbound rule "all traffic" (default) + the subnet has a route <code>0.0.0.0/0 → igw</code>

This is the one reason the backend sits in a *public* subnet in this lab: with no NAT Gateway, a private subnet has no path out, and `dnf install nodejs` / `npm install` would hang. The inbound firewall still blocks everything except the ALB and SSH.

```
# Route table as seen from the instance: the default route goes via the subnet gateway → IGW
ip route
# Prove outbound works
curl -sI https://registry.npmjs.org/ | head -1
# The first-boot log shows the downloads that happened
sudo tail -40 /var/log/user-data.log
```

Hop 6 — Admin → any EC2 (SSH, port 22)

From → to	your laptop → an instance's public IP
Port / protocol	22 / SSH, key-based auth only
Allowed by	<code>bitopi-backend-sg</code> / <code>bitopi-db-sg</code> inbound: SSH 22

```
# Windows: fix the key's permissions once (see scripts/fix-ssh-key-permissions-windows.ps1),
then:
ssh -i bitopi-key.pem ec2-user@<public-ip>

# Copy a file to a server / back
scp -i bitopi-key.pem myfile.txt ec2-user@<public-ip>:/tmp/
scp -i bitopi-key.pem ec2-user@<public-ip>:/var/log/user-data.log .
```

Backend servers **can** SSH to each other too (same key, same security group would need a rule for 22 from `bitopi-backend-sg`) — but nothing in this design needs it, and adding it would only widen the attack surface.

Hop 7 — The control plane: CloudWatch → Auto Scaling Group → EC2

This hop has **no ports and no security groups** — it is AWS services talking to AWS services on your behalf.

1. Every instance's hypervisor reports **CPUUtilization** to CloudWatch (no agent needed for this metric).
2. The target-tracking policy created two alarms: `TargetTracking-bitopi-app-asg-AlarmHigh-...` and `...-AlarmLow-...`.
3. When the group average crosses 30 %, the alarm enters `ALARM` and **calls the Auto Scaling Group**.
4. The ASG computes the new desired capacity and calls **EC2 RunInstances** with the Launch Template (or terminates the instance it selects).
5. The ASG **registers / deregisters** the instance with the Target Group, so the ALB starts / stops sending it traffic. On termination it waits for **connection draining** first.

You can watch all of this in the ASG **Activity** tab — every line quotes the alarm name and the reason.

Quick reference: who can talk to whom

From \ To	ALB :80	Backend :3000	Backend :22	DB :3306	DB :22
Internet	✓	✗	✓ (key required)	✗	✓ (key required)
ALB	—	✓	✗	✗	✗
Backend EC2	✓ (via public DNS)	✗ (not needed)	✗ (no rule)	✓	✗
Database EC2	✓ (via public DNS)	✗	✗	—	—

Every ✓ is one inbound rule in one of the three security groups. Every reply travels back through statefulness. That table *is* the architecture.

4. Part 2 — CI/CD Pipeline for the Auto Scaling Backend

Student: Joyanta Sarker Joy · **Course:** Mastering DevOps, Batch 13 · **Region:** Asia Pacific (Mumbai) `ap-south-1` **Repository:** `JsJoyBitopibd/-AWS-Auto-Scaling-CI-CD-for-3-Tier-Application-Part-1-Auto-Scaling-for-3-Tier-Applications`

1. Objective

Build a pipeline that takes a code change from `git push` all the way to the running Auto Scaling fleet with no manual deployment — and prove that (a) a brand-new server launched by the Auto Scaling Group receives the correct application version by itself, and (b) a new version can be released with a Blue/Green switch and rolled back.

2. Why this pipeline looks different from the assignment's wording

The assignment names **AWS CodePipeline + CodeBuild + CodeDeploy**. Those three services each require an **IAM service role**, and CodeDeploy additionally needs an **EC2 instance profile** attached to the Launch Template (`iam:PassRole`). The course account **cannot create IAM roles** (documented in Part 1, §3), and a check of the CodeDeploy and CodeBuild consoles confirmed no pre-existing roles were available to select.

So the pipeline was built with tools the account *can* use, keeping every stage the assignment asks for:


Assignment stage	AWS-native tool it names	What this project uses	Why it is equivalent
Source	CodePipeline source (GitHub)	GitHub repository, branch <code>main</code>	same source, same trigger (push)
Build + Test	CodeBuild + <code>buildspec.yml</code>	GitHub Actions + <code>.github/workflows/ci-cd.yml</code>	hosted runner installs dependencies and runs the test suite; a failing test stops the pipeline
Artifact	S3 artifact bucket	GitHub Release asset <code>bitopi-app.zip</code> (+ workflow artifact)	versioned, immutable, downloadable over HTTPS
Deploy	CodeDeploy agent + <code>appspec.yml</code>	<code>deploy.sh</code> + a <code>systemd</code> timer on every instance, installed by the Launch Template	same "agent on the box pulls the bundle and restarts the app" model as CodeDeploy — but pull -based, so no AWS credentials are needed anywhere
Blue/Green	CodeDeploy Blue/Green	second Target Group + second Auto Scaling Group + ALB listener switch	this is literally what CodeDeploy does under the hood for EC2 Blue/Green

The result needs **zero IAM roles, zero instance profiles, zero access keys**: nothing ever pushes *into* AWS. Servers reach *out* to GitHub over HTTPS (port 443), which the default security-group egress already allows.

3. Architecture

Part 2 — CI/CD Pipeline for the Auto Scaling Backend (role-free design)

GitHub Actions → GitHub Release → pull-based deploy onto the Auto Scaling fleet · Blue/Green via ALB listener switch | Joyanta Sarker Joy · Mastering DevOps (Batch 13)



The flow, in one paragraph. A developer bumps version in `app/package.json` and pushes to `main`. GitHub Actions checks out the code, installs dependencies, runs the tests, zips the app together with a `VERSION` file and the commit hash, and publishes that zip as a **GitHub Release** tagged `v<version>`. On every backend server a `systemd` timer runs `deploy.sh` every 60 seconds; the script downloads the release for its track, compares the `VERSION` inside with the one running, and — only if different — unzips, `npm install`s and restarts the app. Because the Launch Template runs the same script **at first boot**, any server the Auto Scaling Group creates arrives on the current version before its first health check.

Two release types make Blue/Green possible:

release_type	GitHub marks it	GitHub "latest"	Who deploys it
stable (default on push)	normal release	yes	the Blue fleet — <code>DEPLOY_TRACK=latest</code>
candidate (manual run)	pre-release	no	only a fleet pinned to that tag — Green , <code>DEPLOY_TRACK=v1.2.0</code>

4. Repository contents used by the pipeline

File	Role
app/app.js	Express app: / (page), /health (liveness), /version (what is deployed). Reads its version from package.json, its colour from ENV_COLOR
app/package.json	version → release tag; npm test → node --test test/
app/test/health.test.js	3 tests: /health is 200 + healthy; /version matches package.json; / renders even with no database
.github/workflows/ci-cd.yml	the pipeline (Source → Build → Test → Package → Release)
scripts/deploy.sh	the pull-based deployer that runs on every instance
scripts/app-server-userdata-v2-blue.sh	Launch Template v2 — installs the deployer, ENV_COLOR=blue, DEPLOY_TRACK=latest
scripts/app-server-userdata-v3-green.sh	Launch Template v3 — same, but ENV_COLOR=green, DEPLOY_TRACK=v1.2.0

5. Build steps and evidence

Step 1 — The pipeline: Source → Build → Test → Artifact → Release

What: .github/workflows/ci-cd.yml, triggered by any push to main that touches app/, or run manually with a release_type input.

Why each stage exists:

Step in the workflow	Command	Purpose
Checkout source	actions/checkout@v4	get the code that was pushed
Set up Node.js 20	actions/setup-node@v4	same runtime the servers use
Install dependencies	npm install	build the app
Run application tests	npm test	gate — any failure stops the run; nothing is released
Read version	node -p "require('./package.json').version"	the version is the release tag
Package deployable artifact	zip bitopi-app.zip app.js package.json VERSION COMMIT	the thing that gets deployed; VERSION lets servers detect change, COMMIT ties it to source
Upload artifact	actions/upload-artifact@v4	kept with the run for audit
Publish GitHub Release	gh release create v<version> bitopi-app.zip [--prerelease]	the store the servers pull from

Result — first run, on the first push: the job *Build → Test → Package → Release* succeeded in **13 s**. All **3 tests passed** (# pass 3 / # fail 0). The artifact contained 4 files (5,308 bytes). Release **v1.1.0 (stable)** was published.

Summary

All jobs

Build -> Test -> Package -> Release

Run details

Usage

Workflow file

1 warning and 2 notices

Build -> Test -> Package -> Release
succeeded 1 minute ago in 13s

Search logs

- Set up job 1s
- Checkout source 8s
- Set up Node.js 20 4s
- Install dependencies 3s
- Run application tests 8s
 - Run npm test 4s
 - bitopi-3tier-app@1.1.0 test 6s
 - node --test test/ 7s
 - npm version 1s
 - Subtest: GET /health returns 200 and status "healthy" 1s
 - GET /health returns 200 and status "healthy" 1s
 - duration_ms: 46.26857
 - Subtest: GET /version reports the package.json version 1s
 - GET /version reports the package.json version 1s
 - duration_ms: 7.615181
 - Subtest: GET / renders the page even when the database is unreachable 1s
 - GET / renders the page even when the database is unreachable 1s
 - duration_ms: 19.335727
 - 1..3
 - tests 3
 - suites 0
 - pass 3
 - fail 0
 - cancelled 0
 - skipped 0
 - todo 0
 - duration_ms 369.615627
 - Read version from package.json 8s
 - Package deployable artifact (bitopi-app.zip) 8s

Summary

All jobs

Build -> Test -> Package -> Release

Run details



Usage

Workflow file

Build -> Test -> Package -> Release
succeeded 1 minute ago in 13s

Search logs

- Run application tests 8s
 - Subtest: GET /health returns 200 and status "healthy" 1s
 - GET /health returns 200 and status "healthy" 1s
 - duration_ms: 46.26857
 - Subtest: GET /version reports the package.json version 1s
 - GET /version reports the package.json version 1s
 - duration_ms: 7.615181
 - Subtest: GET / renders the page even when the database is unreachable 1s
 - GET / renders the page even when the database is unreachable 1s
 - duration_ms: 19.335727
 - 1..3
 - tests 3
 - suites 0
 - pass 3
 - fail 0
 - cancelled 0
 - skipped 0
 - todo 0
 - duration_ms 369.615627
- Read version from package.json 8s
- Package deployable artifact (bitopi-app.zip) 8s
- Upload artifact to this workflow run 1s
- Publish GitHub Release v1.1.0 2s
 - Run set -e 1s
 - https://github.com/330n8itoupbd/-aws-auto-scaling-ci-cd-for-3-tier-application-part-1-auto-scaling-for-3-tier-applications/releases/tag/v1.1.0 1s
 - notice: v1.1.0 (stable) -> https://github.com/330n8itoupbd/-aws-auto-scaling-ci-cd-for-3-tier-application-part-1-auto-scaling-for-3-tier-applications/releases/tag/v1.1.0 1s
- Post Set up Node.js 20 8s
- Post Checkout source 8s
- Complete job 8s



Step 2 — The deployer on every server (Launch Template v2)

What: a new Launch Template version whose user-data no longer contains the application. Instead it installs:

- `/opt/deploy/deploy.sh` — the deployer (below)
- `/opt/deploy/deploy.env` — `DEPLOY_TRACK=latest`
- `bitopi-deploy.service` + `bitopi-deploy.timer` — runs the deployer 45 s after boot and every 60 s after that
- `/opt/app/.env` — instance ID, AZ, `ENV_COLOR`, database connection (**not** part of the release zip, so a deploy never overwrites configuration)


...and then runs the deployer **once immediately**, so the app is installed before the ALB's first health check.


The deployer, in plain words ([scripts/deploy.sh](#)):

```
TRACK="${DEPLOY_TRACK:-latest}" # "latest" or a tag such as v1.2.0
URL="../../../releases/latest/download/bitopi-app.zip" # or ../../../releases/download/$TRACK/bitopi-
app.zip
curl -fsSL -o app.zip "$URL" || exit 0 # can't reach GitHub? do nothing, try again
in 60 s
unzip -q -o app.zip -d tmp
NEW=$(cat tmp/VERSION); CUR=$(cat /opt/app/VERSION || echo none)
[ "$NEW" = "$CUR" ] && exit 0 # already running this version → nothing to
do
cp -r tmp/. /opt/app/ && cd /opt/app && npm install --omit=dev
systemctl restart bitopi-app # ~1 s; /health is shallow so the ALB never
notices
echo "deployed $NEW (commit $(cat /opt/app/COMMIT))" >> /var/log/bitopi-deploy.log
```

Why a timer instead of a push: a push-based deploy (CodeDeploy, SSH from a runner) needs something outside the VPC to hold credentials for the servers. A pull needs nothing: the server already has outbound internet, and GitHub Releases are public. It also means **new servers deploy themselves** (Step 4) with no extra machinery.


Rolling the fleet onto v2 — Instance Refresh. The two running servers still had the old v1 recipe, so the Auto Scaling Group's *Instance refresh* replaced them one at a time (rolling strategy, 120 s warm-up) — the ALB kept serving throughout.






Step 3 — Verify the deployed application

After the refresh, the ALB served **v1.1.0** from a brand-new instance (`i-09a53754f7dfa27b`), environment **blue**, database **CONNECTED** — the page footer reads "*Deployed by GitHub Actions → GitHub Release → pull-based deployer*". `/version` confirms it as JSON.







Step 4 — Automatic deployment: push a change, touch nothing in AWS

Test: bump `package.json` to **1.1.1**, add one visible line to the page, `git push`. No AWS console action of any kind.

Time (UTC)	What happened	Where
~07:19	push → Actions run → tests pass → Release v1.1.1 (stable)	GitHub
07:20:36	<code>[latest] deploying 1.1.0 -> 1.1.1</code>	<code>/var/log/bitopi-deploy.log</code> on <code>i-09a53754f7dfa27b</code>
07:20:37	<code>[latest] deployed 1.1.1 (commit 15b3fbb64e...)</code>	same
~07:21	second instance picks it up on its own next tick	—

The commit hash in the server log is the exact commit that was pushed. The page now shows **v1.1.1** and the yellow *"What's new in v1.1.1: this line reached the server through automatic pull-based deployment — no AWS console action was taken."* Deploy-to-live on the whole fleet: **under two minutes from `git push`, one second per server once its timer fires.**



bitopi-app-asg Capacity overview

arn:aws:autoscaling:ap-south-1:448513989308:autoScalingGroup:5fb9be00-6794-4c7f-8fb6-cab96bbb2f0d:autoScalingGroupName/bitopi-app-asg

Desired capacity: 3 Scaling limits: 2 - 4 Desired capacity type: Units (number of instances) Status: Updating capacity

Date created: Fri Sep 25 2026 10:38:46 GMT+0600 (Bangladesh Standard Time)

Details | Integrations | Automatic scaling | Instance management | Instance refresh | Activity | Monitoring | Tags

Activity notifications (0)

Filter notifications

Send to: On instance action

No notifications are currently specified

Create notification

Activity history (12)

Filter activity history

Filter by a date and time range

Status	Description	Cause	Start time	End time
Successful	Launching a new EC2 instance: i-025515fd7253ffa4d	At 2026-09-25T07:14:27Z an instance was launched in response to an instance refresh.	2026 September 25, 01:14:28 PM +06:00	2026 September 25, 01:16:33 PM +06:00
Successful	Terminating EC2 instance: i-0f33782fbae77fa57	At 2026-09-25T07:14:26Z an instance was taken out of service in response to an instance refresh.	2026 September 25, 01:14:27 PM +06:00	2026 September 25, 01:20:09 PM +06:00

bitopi-app-tg

Details

arn:aws:elasticloadbalancing:ap-south-1:448513989308:targetgroup/bitopi-app-tg/76b307f2902dbf04

Target type: Instance Protocol : Port: HTTP: 3000 Protocol version: HTTP1 VPC: vpc-0ab019c806b84a938

IP address type: IPv4 Load balancer: bitopi-alb

3 Total targets 3 Healthy 0 Unhealthy 0 Unused 0 Initial 0 Draining

0 Anomalous

Distribution of targets by Availability Zone (AZ)

Select values in this table to see corresponding filters applied to the Registered targets table below.

Targets | Monitoring | Health checks | Attributes | Tags

Registered targets (3)

Anomaly mitigation: Not applicable

Deregister | Register targets

Target groups route requests to individual registered targets using the protocol and port number specified. Health checks are performed on all registered targets according to the target group's health check settings. Anomaly detection is automatically applied to HTTP/HTTPS target groups with at least 3 healthy targets.

Filter targets

Instance ID	Name	Port	Zone	Health status	Health status details	Administrative o...	Override details	Launch...	Anomaly
i-066f9aa10f6c4556a	bitopi-app-ser...	3000	ap-south-1b (a...	Healthy	-	No override	No override is curren...	Septembe...	Normal
i-025515fd7253ffa4d	bitopi-app-ser...	3000	ap-south-1a (a...	Healthy	-	No override	No override is curren...	Septembe...	Normal
i-09a537547dfda277b	bitopi-app-ser...	3000	ap-south-1b (a...	Healthy	-	No override	No override is curren...	Septembe...	Normal

```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

PS C:\WINDOWS\System32> Copy-Item E:\DevSecOps\Assignments\9_Module_Assignment\SS\P2-14_new-instance-serves-v1.1.1.png E:\DevSecOps\Assignments\9_Module_Assignment\SS\P2-10_app-v1.1.1-auto-deployed.png -Force
PS C:\WINDOWS\System32> 1..12 | ForEach-Object { curl.exe -s http://bitopi-alb-1957667478.ap-south-1.elb.amazonaws.com/version; "" }
{"version":"1.1.1","env":"blue","instance":"i-025515fd7253ffa4d"}
{"version":"1.1.1","env":"blue","instance":"i-09a53754f7dfda27b"}
{"version":"1.1.1","env":"blue","instance":"i-025515fd7253ffa4d"}
{"version":"1.1.1","env":"blue","instance":"i-09a53754f7dfda27b"}
{"version":"1.1.1","env":"blue","instance":"i-09a53754f7dfda27b"}
{"version":"1.1.1","env":"blue","instance":"i-025515fd7253ffa4d"}
{"version":"1.1.1","env":"blue","instance":"i-09a53754f7dfda27b"}
{"version":"1.1.1","env":"blue","instance":"i-025515fd7253ffa4d"}
{"version":"1.1.1","env":"blue","instance":"i-09a53754f7dfda27b"}
{"version":"1.1.1","env":"blue","instance":"i-025515fd7253ffa4d"}
{"version":"1.1.1","env":"blue","instance":"i-09a53754f7dfda27b"}
{"version":"1.1.1","env":"blue","instance":"i-025515fd7253ffa4d"}
{"version":"1.1.1","env":"blue","instance":"i-09a53754f7dfda27b"}
PS C:\WINDOWS\System32> |
```

```
ec2-user@ip-10-0-24-151:~$ npm notice
2026-09-25T07:13:00+00:00 [latest] deployed 1.1.0 (commit 86f233769d51a949a23012f6fa703f7959af3f05)
2026-09-25T07:20:36+00:00 [latest] deploying 1.1.0 -> 1.1.1

up to date in 516ms
2026-09-25T07:20:37+00:00 [latest] deployed 1.1.1 (commit 15b3fbb64e593651c59f812d92de0f2c98aa3616)
---
1.1.1

2026-09-25 07:12:17
[ec2-user@ip-10-0-24-151 ~]$ sudo cat /var/log/bitopi-deploy.log; echo "----"; cat /opt/app/VERSION; echo; uptime -s
2026-09-25T07:12:53+00:00 [latest] deploying none -> 1.1.0

added 79 packages in 7s
npm notice
npm notice New major version of npm available! 10.8.2 -> 12.1.0
npm notice Changelog: https://github.com/npm/cli/releases/tag/v12.1.0
npm notice To update run: npm install -g npm@12.1.0
npm notice
2026-09-25T07:13:00+00:00 [latest] deployed 1.1.0 (commit 86f233769d51a949a23012f6fa703f7959af3f05)
2026-09-25T07:20:36+00:00 [latest] deploying 1.1.0 -> 1.1.1

up to date in 516ms
2026-09-25T07:20:37+00:00 [latest] deployed 1.1.1 (commit 15b3fbb64e593651c59f812d92de0f2c98aa3616)
---
1.1.1

2026-09-25 07:12:17
[ec2-user@ip-10-0-24-151 ~]$ |
```

This is the property that ties the pipeline to the Auto Scaling Group: **scale-out at 3 a.m. and self-healing after a crash both produce servers on the right version**, because the template only knows *how to fetch*, never *what*.

Step 6 — Blue/Green deployment and rollback (executed)

Design.

```

┌ listener :80 default action ─┐ (production – exactly one colour
at a time)
Internet ──> bitopi-alb
└ listener :8080 ─ test only ─┘ ──> always GREEN (used to test
before the switch)
└─> TG bitopi-app-tg ──> ASG bitopi-app-asg (BLUE, LT
v2, DEPLOY_TRACK=latest → v1.1.1)
└─> TG bitopi-app-tg-green ──> ASG bitopi-app-asg-green (GREEN, LT
v3, DEPLOY_TRACK=v1.2.0)
both fleets share the same MySQL
server (10.0.13.216)

```

Element	Blue	Green
Target group	bitopi-app-tg	bitopi-app-tg-green (identical settings, /health)
Auto Scaling Group	bitopi-app-asg (min 2 / max 4)	bitopi-app-asg-green (min 1 / max 2 for the lab)
Launch Template	bitopi-app-lt v2 (default, and the version the blue ASG is pinned to)	bitopi-app-lt v3 — app-server-userdata-v3-green.sh
Release track	latest (stable releases only)	pinned v1.2.0 (a pre-release)
Page badge	BLUE ENVIRONMENT, blue accent	GREEN ENVIRONMENT, green accent

6.0 Resuming on a partly reclaimed account (and one real mistake caught)

The course account reclaims EC2 servers after a while, so the work resumed with the read-only **Terraform audit** (§9). It reported everything **PRESENT** — VPC, subnets, security groups, key pair, launch template v2, both target groups, ALB, blue ASG — except the **database server**, which had been reclaimed. The blue app servers had survived because the Auto Scaling Group replaces anything that disappears.

```

main.tf
1 # -----
2 # Bitopi 3-Tier assignment - Terraform "is it still there?" audit
3
4 -----
5
6 PS E:\DevSecOps\Assignments\9_Module_Assignment\terraform> terraform apply -auto-approve
7
8 alb_url = "http://bitopi-alb-1957667478.ap-south-1.elb.amazonaws.com/"
9 audit_summary = <<EOT
10
11 ----- BITOPI 3-TIER AUDIT (ap-south-1) -----
12 VPC bitopi-vpc ..... PRESENT
13 4 subnets bitopi-subnet-* ..... PRESENT (found 4)
14 3 security groups ..... PRESENT (found 3: bitopi-backend-sg, bitopi-db-sg, bitopi-alb-sg)
15 DB server bitopi-db-server ..... MISSING
16 ALB bitopi-alb ..... PRESENT bitopi-alb-1957667478.ap-south-1.elb.amazonaws.com
17 listener :80 currently forwards to ..... bitopi-app-tg
18 ASG bitopi-app-asg (blue) ..... PRESENT min 2 / max 4, LT bitopi-app-lt v2
19 ASG bitopi-app-asg-green (green) ..... MISSING (only needed for the Blue/Green step)
20 key pair / launch template / target groups : see the WARNINGS above
21 (no warning = PRESENT)
22
23 -----
24 Everything marked MISSING (or named in a warning) must be rebuilt by
25 hand in the console - follow docs/part1-auto-scaling.md for Part 1 and
26 docs/part2-cicd.md section 7 for the Blue/Green steps.
27
28 EOT
29 blue_asg = {
30   "desired" = 2
31   "health_check_type" = "ELB"
32   "launch_template_version" = "2"
33   "max" = 4
34   "min" = 2
35   "target_groups" = toset([
36     "arn:aws:elasticloadbalancing:ap-south-1:448513989308:targetgroup/bitopi-app-tg/f6b307f2902dbf04",
37   ])
38 }
39 listener_forwards_to = "bitopi-app-tg"
40 public_subnet_ids = {
41   "bitopi-subnet-public1-ap-south-1a" = "subnet-053a27c29e6c271ad"
42   "bitopi-subnet-public2-ap-south-1b" = "subnet-0d4371e7ff2dcf8"
43 }
44 security_group_ids = {
45   "bitopi-alb-sg" = "sg-04c8272fbee0e769"
46   "bitopi-backend-sg" = "sg-0029472c273fdd274"
47 }

```

The audit also surfaced a mistake worth documenting. The candidate build (workflow run #3 on branch `release/v1.2.0`) had been run with the default `release_type = stable`, so GitHub published **v1.2.0 as a normal release marked Latest**. The blue fleet follows `latest`, so it had quietly installed v1.2.0 — exactly the failure mode safeguard (b) below is meant to prevent (the release page showed 2,710 downloads of the v1.2.0 zip: every blue server checking once a minute). The fix needed no AWS action: v1.2.0 was edited to **Pre-release**, v1.1.1 became *Latest* again, and within a minute every blue server's deployer saw a different `VERSION` and rolled itself back to 1.1.1.

The screenshot shows a GitHub Actions workflow run for the repository `IsJoyBitopibd/-AWS-Auto-Scaling-CI-CD-for-3-Tier-Application-Part-1-Auto-Scaling-for-3-Tier-Applications`. The workflow is named `CI/CD - build, test, release #3` and is currently in a **Success** state. It was manually triggered by user `JJoyBitopibd` on the `release/v1.2.0` branch. The total duration of the run is 22 seconds, and it produced one artifact.

The workflow steps are summarized as `Build → Test → Package → Release`. The run details show that the workflow file was executed successfully in 16 seconds.

Annotations: A warning indicates that the `node.js 20` is deprecated and the following actions target `node.js 20` but are being forced to run on `node.js 24`: `actions/checkout@v4`, `actions/setup-node@v4`, and `actions/upload-artifact@v4`. For more information, see <https://github.blog/changelog/2025-09-19-deprecati...>.

Released: v1.2.0 (stable) -> <https://github.com/IsJoyBitopibd/-AWS-Auto-Scaling-CI-CD-for-3-Tier-Application-Part-1-Auto-Scaling-for-3-Tier-Applications/releases/tag/v1.2.0>

Artifacts: Produced during runtime. The artifact table shows:

Name	Size	Digest
bitopi-app-v1.2.0	2.86 KB	sha256:288453280c368f25489c886251143238fc84898ee2b481f84423983d159d9fa

The screenshot shows the GitHub Releases page for the repository `IsJoyBitopibd/-AWS-Auto-Scaling-CI-CD-for-3-Tier-Application-Part-1-Auto-Scaling-for-3-Tier-Applications`. The page displays a list of releases, with **v1.2.0** currently selected and marked as **Pre-release**.



Release list:

- v1.2.0** (Pre-release): Released this 2 days ago. Built from `7f8859` by GitHub Actions run #3 (stable). Assets: 3.
- v1.1.1** (Latest): Released this 2 days ago. Built from `15b3fbb` by GitHub Actions run #2 (stable). Assets: 3.
- v1.1.0**

Assets for v1.1.1:

Asset	Size	Uploaded
bitopi-app.zip	2.94 KB	2 days ago
Source code (zip)		2 days ago
Source code (tar.gz)		2 days ago

Database rebuilt with the same private IP. `bitopi-db-server` was relaunched from the same user-data with *Primary IP* set to `10.0.13.216` (inside `bitopi-subnet-public1-ap-south-1a`, `10.0.0.0/20`). Because the address is the one baked into the launch templates as `DB_HOST`, nothing else had to change: the running blue servers reconnected on their next request.



6.1 Deploy the new version to Green

1. **Candidate release** — v1.2.0 (P2-16/P2-17 above) exists only as a pre-release, so only a fleet pinned to the tag installs it.
2. **Green target group** bitopi-app-tg-green — HTTP 3000, health check /health, same settings as blue.
3. **Launch Template v3** — created from v2 with the green user-data (ENV_COLOR=green , DEPLOY_TRACK=v1.2.0 , DB_HOST=10.0.13.216). **Default stays 2**, and the blue ASG is pinned to version 2, so blue cannot pick up v3.
4. **Green ASG** bitopi-app-asg-green — LT **version 3**, both public subnets, target group bitopi-app-tg-green , ELB health checks, min 1 / desired 1 / max 2. Its instance i-02334b27eb02d41fe booted and its deployer fetched **v1.2.0** on first boot.

ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#TargetGroup:targetGroupArn=arn:aws:elasticloadbalancing:ap-south-1:448513989308:targetgroup/bitopi-app-tg-green/889c2dd7d247d442...

EC2 > Target groups > bitopi-app-tg-green

bitopi-app-tg-green

Details

arn:aws:elasticloadbalancing:ap-south-1:448513989308:targetgroup/bitopi-app-tg-green/889c2dd7d247d442

Target type
Instance

Protocol : Port
HTTP: 3000

Protocol version
HTTP1

VPC
vpc-0ab019c806b84a938

IP address type
IPv4

Load balancer
None associated

0 Total targets

0 Healthy

0 Unhealthy

0 Unused

0 Initial

0 Draining

0 Anomalous

Targets | Monitoring | Health checks | Attributes | Tags

Registered targets (0) [Info](#)

Target groups route requests to individual registered targets using the protocol and port number specified. Health checks are performed on all registered targets according to the target group's health check settings. Anomaly detection is automatically applied to HTTP/HTTPS target groups with at least 3 healthy targets.

[Anomaly mitigation: Not applicable](#) [Deregister](#) [Register targets](#)

Filter targets

Instance ID	Name	Port	Zone	Health status	Health status details	Administrative o...	Override details	Launch time
No registered targets								

You have not registered targets to this group yet

[Register targets](#)

ap-south-1.console.aws.amazon.com/ec2/home?region=ap-south-1#LaunchTemplate:launchTemplateId=lt-072cb16544d51e361

EC2 > Launch templates > bitopi-app-lt

bitopi-app-lt (lt-072cb16544d51e361)

Launch template details

Launch template ID
lt-072cb16544d51e361

Launch template name
bitopi-app-lt


Default version
2

Owner
arn:aws:iam::448513989308:user/joyanta15-2934@diu.edu.bd

Details | **Versions** | Template tags


Versions (3)


Version	Default Version	Description	Creation Time	Created By	AMI ID
3	No	v3 - GREEN, pinned to candidate v1.2.0	2026-09-27T09:22:32.000Z	arn:aws:iam::448513989308:user/joyanta15-2934@diu.edu.bd	ami-066c4849e6b3a1e3d
2	Yes	v2 - pull-based deployer from GitHub Releases (blue)	2026-09-25T07:09:09.000Z	arn:aws:iam::448513989308:user/joyanta15-2934@diu.edu.bd	ami-066c4849e6b3a1e3d
1	No	v1 - Node app + DB wiring	2026-09-25T04:20:15.000Z	arn:aws:iam::448513989308:user/joyanta15-2934@diu.edu.bd	ami-066c4849e6b3a1e3d



6.2 Health checks before any traffic moves


A target group is only health-checked when a listener sends traffic to it, so right after creation the green target showed **"Unused — target group is not configured to receive traffic from the load balancer"**. To test green without exposing it to users, a **test listener** was added: **bitopi-alb HTTP:8080 → bitopi-app-tg-green** (plus an inbound rule for TCP 8080 on bitopi-alb-sg). Within a minute the green target was **healthy**, and **http://<alb>:8080/** showed **v1.2.0, GREEN ENVIRONMENT, database CONNECTED** — while the production URL on port 80 still showed v1.1.1 blue.







6.3 Switch production traffic

bitopi-alb → listener **HTTP:80** → *Edit* → default action **forward to bitopi-app-tg-green (100 %)**. The production URL immediately served **v1.2.0 GREEN**. Blue was left running, untouched. The Terraform audit, run independently against the AWS API, confirmed `listener_forwards_to = "bitopi-app-tg-green"` and the green ASG on launch template v3.





6.4 Rollback test



A PowerShell loop requested `/version` through the ALB once a second while the listener was edited back to `bitopi-app-tg`:



```

$u = "http://bitopi-alb-1957667478.ap-south-1.elb.amazonaws.com/version"
while ($true) { try { $r = Invoke-RestMethod $u -TimeoutSec 3
  "{0} v{1} {2} {3}" -f (Get-Date -Format HH:mm:ss), $r.version, $r.env, $r.instance }
  catch { "{0} ERROR {1}" -f (Get-Date -Format HH:mm:ss), $_.Exception.Message }; Start-Sleep 1
}
  
```

Time (UTC+6)	Response
16:41:16	v1.2.0 green i-02334b27eb02d41fe — last green answer
16:41:17	v1.1.1 blue i-0bd8f786ce68ed0f9 — first blue answer
16:41:18 →	alternating i-0bd8f786ce68ed0f9 / i-0aec18b50bb27272a , all v1.1.1 blue

Result: the rollback took effect between two consecutive one-second samples, and not a single request failed (no ERROR line). Nothing was redeployed — blue had never been changed, which is the whole point of Blue/Green compared with a rolling update, where a rollback means deploying the old version all over again. Terraform again confirmed `listener_forwards_to = "bitopi-app-tg"`.





Two safeguards, both proven necessary. (a) Blue's ASG references a **fixed** Launch Template version (2), not *Latest* — otherwise creating v3 would have made blue launch green servers. (b) A candidate must be published as a **pre-release** (`release_type = candidate`); \$6.0 shows what happens when it is not: the blue fleet deploys it automatically, because that is exactly what the pipeline is built to do.

6. Submission checklist → evidence

Required item	Where
CI/CD architecture diagram	diagrams/part2_cicd_architecture.png
Pipeline configuration	.github/workflows/ci-cd.yml ; P2-01, P2-02
Build configuration	workflow steps <i>Install dependencies</i> , <i>Run application tests</i> ; P2-01
Deployment configuration	scripts/deploy.sh , scripts/app-server-userdata-v2-blue.sh ; P2-05, P2-06
Successful deployment evidence	P2-07, P2-08, P2-09, P2-10, P2-11
New Auto Scaling instance deployment evidence	P2-12, P2-13, P2-14, P2-15
Blue/Green deployment	§5 Step 6; app-server-userdata-v3-green.sh ; P2-16 ... P2-26, TF-02
Rollback test evidence	§6.4; P2-27, P2-28, P2-29, TF-03
Stack verification (Terraform audit)	terraform/ ; §9; TF-01, TF-02, TF-03
Final documentation	this file + README.md + docs/how-servers-communicate.md + docs/Bitopi_AWS_AutoScaling_CICD_Assignment.pdf

7. Runbook — re-running Blue/Green on a rebuilt stack (~20 min)

Use this if the account reclaims the servers again.

1. **Audit.** terraform apply -auto-approve in terraform/ (see [terraform/README.md](#)). Rebuild only what is MISSING, in Part 1 order.
2. **DB server with the same IP.** If bitopi-db-server is missing, relaunch it with *Advanced network configuration* → *Primary IP* = 10.0.13.216, so no launch template needs editing.
3. **Candidate release.** On release/v1.2.0, run the workflow with release_type = **candidate**. Check *Releases: v1.1.1 Latest, v1.2.0 Pre-release*. If v1.2.0 shows *Latest*, edit it → tick *Set as a pre-release* → *Update release*.
4. **Protect blue.** Blue ASG must use Launch Template version **2** (a fixed number, not *Latest*).
5. **Green TG** bitopi-app-tg-green (HTTP 3000, /health), **LT v3** from v2 with the green user-data (not default), **green ASG** on version 3.
6. **Test listener.** bitopi-alb-sg inbound TCP 8080; ALB listener HTTP:8080 → green TG. Wait for *healthy*, browse :8080.
7. **Switch.** Listener HTTP:80 → bitopi-app-tg-green. Verify page + terraform output listener_forwards_to.
8. **Rollback.** Listener HTTP:80 → bitopi-app-tg, with the /version monitor loop running. Verify again.
9. **Clean up the green side** (§8).

8. Cleanup checklist (avoid charges)

Delete in this order (each step waits for the previous to finish):

1. Auto Scaling Groups `bitopi-app-asg-green` and `bitopi-app-asg` — delete (this terminates their instances)
2. ALB listener HTTP:8080 and the TCP 8080 inbound rule on `bitopi-alb-sg`; then load balancer `bitopi-alb`; then target groups `bitopi-app-tg`, `bitopi-app-tg-green`
3. EC2 instance `bitopi-db-server`
4. Launch template `bitopi-app-lt` (all versions)
5. Key pair `bitopi-key`
6. Security groups `bitopi-db-sg`, `bitopi-backend-sg`, `bitopi-alb-sg` (in that order — each is referenced by the previous)
7. VPC `bitopi-vpc` (deletes its subnets, route tables and internet gateway)

The GitHub repository, releases and workflow cost nothing and stay.

9. Verifying the whole stack at any time — the Terraform audit

Because the course account reclaims servers, "is everything still working?" is a question that comes up every time work resumes. Instead of clicking through six console pages, the repository carries a **read-only Terraform program** in `terraform/` that answers it in one command.

What it is. A Terraform configuration with *no* `resource` blocks — only lookups (`data` sources) and `check` blocks. `terraform apply` therefore cannot create, change or delete anything; it can only describe. It needs the same `Describe*` permissions the console already uses.

What it reports. One line per resource — VPC, the 4 subnets, the 3 security groups, key pair, DB server, launch template (and whether version 2 with the deployer exists), target group (port 3000, `/health`), ALB, Auto Scaling Group, and the optional green pair — as `PRESENT` or `MISSING`, plus the values needed to continue: the DB server's private IP (= `DB_HOST` in the user-data), the ALB URL, the ASG's min/desired/max and launch-template version, and **which target group the :80 listener is forwarding to** (`bitopi-app-tg` = blue, `bitopi-app-tg-green` = green). That last value is direct evidence for the Blue/Green switch and the rollback.

How to run it (PowerShell, once Terraform is installed and an access key is in the environment — full step-by-step in `terraform/README.md`):

```
cd E:\DevSecOps\Assignments\9_Module_Assignment\terraform
terraform init
terraform apply -auto-approve
terraform output audit_summary
```

How to read it. Warnings named `check "<resource>"` mean that resource is missing or mis-configured (the message names the Part/Step that rebuilds it); the `audit_summary` box lists `PRESENT/MISSING` for everything else. No warnings and every line `PRESENT` = the stack is complete. The only warning that is *expected* before the Blue/Green step is `target_group_green_optional`.

Where it fits in the workflow.

When	Command	Expected
Resuming work after the servers were reclaimed	<code>terraform apply -auto-approve</code>	tells you exactly what to rebuild by hand (§7)
After rebuilding	same	everything <code>PRESENT</code> , only the green warning left
After the Blue/Green switch	<code>terraform output listener_forwards_to</code>	<code>bitopi-app-tg-green</code>
After the rollback	same	<code>bitopi-app-tg</code>

Evidence: `SS/TF-01-terraform-audit-before.png` (state found on resuming — only the DB server missing), `SS/TF-02-terraform-after-switch.png` (everything present, listener on green), `SS/TF-03-terraform-after-rollback.png` (listener back on blue).

Practical notes from real use. `terraform output` does **not** contact AWS — it prints what the last *successful* `apply` saved in `terraform.tfstate`. Once, the office DNS server briefly failed (`lookup sts.ap-south-1.amazonaws.com: no such host`); `apply` aborted and `output` still showed the previous answer. Always check that `apply` ended with `Apply complete!` before trusting the outputs (`ipconfig /flushdns` and a retry fixed it).

5. Appendix — scripts and pipeline

db-server-userdata.sh

```
#!/bin/bash
# =====
# Bitopi 3-Tier Assignment - DATABASE SERVER user-data
# Runs automatically on first boot of the DB EC2 instance.
# Installs MariaDB (MySQL-compatible), creates the app database,
# an app user, and a seed table. No manual steps needed.
# =====
exec >>(tee /var/log/user-data.log) 2>&1
set -x

# 1) Install and start MariaDB (MySQL-compatible server)
dnf install -y mariadb105-server
systemctl enable --now mariadb

# 2) Make MySQL listen on all interfaces so the app servers can reach it
cat >/etc/my.cnf.d/99-app.cnf <<'EOF'
[mysqld]
bind-address=0.0.0.0
EOF
systemctl restart mariadb

# 3) Create the application database, a remote-capable user, and seed data
mysql <<'SQL'
CREATE DATABASE IF NOT EXISTS appdb;
CREATE USER IF NOT EXISTS 'appuser'@'%' IDENTIFIED BY 'Bitopi#Db2026';
GRANT ALL PRIVILEGES ON appdb.* TO 'appuser'@'%';
FLUSH PRIVILEGES;
USE appdb;
CREATE TABLE IF NOT EXISTS app_info (id INT PRIMARY KEY, info VARCHAR(255));
INSERT IGNORE INTO app_info (id, info)
VALUES (1, 'Bitopi 3-Tier App - data served live from MySQL on EC2');
CREATE TABLE IF NOT EXISTS visits (
  id INT AUTO_INCREMENT PRIMARY KEY,
  hostname VARCHAR(255),
  visited_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
SQL

echo "==== Bitopi DB setup complete ====="
```

app-server-userdata-v2-blue.sh

```
#!/bin/bash
# =====
# Bitopi 3-Tier - APP SERVER user-data v2 (BLUE environment)
# Part 2: the app is no longer embedded here. Every server installs a
# pull-based deployer that fetches the app from GitHub Releases:
#   - at boot -> a brand-new Auto Scaling instance gets the current
#               release before its first health check
#   - every 60s -> a new release rolls out with no AWS action at all
# DEPLOY_TRACK: latest = newest stable release | vX.Y.Z = pinned tag
# =====
exec >>(tee /var/log/user-data.log) 2>&1
set -x

# 1) Runtime + tools
dnf install -y nodejs unzip

# 2) Who am I? (IMDSv2)
TOKEN=$(curl -s -X PUT "http://169.254.169.254/latest/api/token" -H "X-aws-ec2-metadata-token-ttl-seconds: 300")
IID=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" http://169.254.169.254/latest/meta-data/instance-id)
AZ=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" http://169.254.169.254/latest/meta-data/placement/availability-zone)

mkdir -p /opt/app /opt/deploy

# 3) App configuration (never inside the release zip, so deploys keep it)
cat >/opt/app/.env <<ENV
PORT=3000
INSTANCE_ID=$IID
AZ=$AZ
ENV_COLOR=blue
DB_HOST=10.0.13.216
DB_USER=appuser
DB_PASS=Bitopi#Db2026
DB_NAME=appdb
DB_PORT=3306
ENV

# 4) Which release track this environment follows
cat >/opt/deploy/deploy.env <<ENV
DEPLOY_TRACK=latest
ENV

# 5) The deployer (identical to scripts/deploy.sh in the repo)
cat >/opt/deploy/deploy.sh <<'DEPLOY'
#!/bin/bash
# =====
# Bitopi pull-based deployer - /opt/deploy/deploy.sh
# Runs at boot and every 60 s (systemd timer). Fetches the release for
# its track from GitHub, and if the VERSION differs from what is running,
# installs it and restarts the app. No AWS credentials required.
#
# DEPLOY_TRACK=latest -> newest stable GitHub Release (BLUE fleet)
# DEPLOY_TRACK=v1.2.0 -> that exact tag, even a pre-release (GREEN fleet)
# =====
set -uo pipefail
REPO="JsJoyBitopibd/-AWS-Auto-Scaling-CI-CD-for-3-Tier-Application-Part-1-Auto-Scaling-for-3-Tier-Applications"
TRACK="${DEPLOY_TRACK:-latest}"
APP_DIR=/opt/app
LOG=/var/log/bitopi-deploy.log
exec >>"$LOG" 2>&1

if [ "$TRACK" = "latest" ]; then
    URL="https://github.com/$REPO/releases/latest/download/bitopi-app.zip"
```

```

else
    URL="https://github.com/$REPO/releases/download/$TRACK/bitopi-app.zip"
fi

TMP=$(mktemp -d)
if ! curl -fsSL --retry 3 -o "$TMP/app.zip" "$URL"; then
    echo "$(date -Is) [$TRACK] download failed: $URL"; rm -rf "$TMP"; exit 0
fi
unzip -q -o "$TMP/app.zip" -d "$TMP/app"
NEW=$(cat "$TMP/app/VERSION")
CUR=$(cat "$APP_DIR/VERSION" 2>/dev/null || echo "none")

if [ "$NEW" = "$CUR" ]; then rm -rf "$TMP"; exit 0; fi # nothing to do

echo "$(date -Is) [$TRACK] deploying $CUR -> $NEW"
mkdir -p "$APP_DIR"
cp -r "$TMP/app/." "$APP_DIR/" # .env is not in the zip, so it is preserved
cd "$APP_DIR" && npm install --omit=dev --no-audit --no-fund
systemctl restart bitopi-app
echo "$(date -Is) [$TRACK] deployed $NEW (commit $(cat "$APP_DIR/COMMIT" 2>/dev/null))"
rm -rf "$TMP"
DEPLOY
chmod +x /opt/deploy/deploy.sh

# 6) systemd: the app service, and the deploy timer that checks every minute
cat >/etc/systemd/system/bitopi-app.service <<'SVC'
[Unit]
Description=Bitopi 3-Tier App
After=network-online.target
[Service]
EnvironmentFile=/opt/app/.env
WorkingDirectory=/opt/app
ExecStart=/usr/bin/node /opt/app/app.js
Restart=always
User=root
[Install]
WantedBy=multi-user.target
SVC

cat >/etc/systemd/system/bitopi-deploy.service <<'SVC'
[Unit]
Description=Bitopi pull-based deployer (fetch release, install if new)
[Service]
Type=oneshot
EnvironmentFile=/opt/deploy/deploy.env
ExecStart=/opt/deploy/deploy.sh
SVC

cat >/etc/systemd/system/bitopi-deploy.timer <<'TMR'
[Unit]
Description=Check GitHub Releases for a new app version every 60 seconds
[Timer]
OnBootSec=45
OnUnitActiveSec=60
AccuracySec=5
[Install]
WantedBy=timers.target
TMR

systemctl daemon-reload

# 7) First deploy NOW, so the app is running before the ALB health checks start
set -a; . /opt/deploy/deploy.env; set +a
/opt/deploy/deploy.sh

systemctl enable --now bitopi-app

```

```
systemctl enable --now bitopi-deploy.timer
echo "==== Bitopi app server (blue, track=latest) ready ====="
```

deploy.sh

```
#!/bin/bash
# =====
# Bitopi pull-based deployer - /opt/deploy/deploy.sh
# Runs at boot and every 60 s (systemd timer). Fetches the release for
# its track from GitHub, and if the VERSION differs from what is running,
# installs it and restarts the app. No AWS credentials required.
#
# DEPLOY_TRACK=latest -> newest stable GitHub Release (BLUE fleet)
# DEPLOY_TRACK=v1.2.0 -> that exact tag, even a pre-release (GREEN fleet)
# =====
set -uo pipefail
REPO="JsJoyBitopibd/-AWS-Auto-Scaling-CI-CD-for-3-Tier-Application-Part-1-Auto-Scaling-for-3-
Tier-Applications"
TRACK="${DEPLOY_TRACK:-latest}"
APP_DIR=/opt/app
LOG=/var/log/bitopi-deploy.log
exec >>"$LOG" 2>&1

if [ "$TRACK" = "latest" ]; then
    URL="https://github.com/$REPO/releases/latest/download/bitopi-app.zip"
else
    URL="https://github.com/$REPO/releases/download/$TRACK/bitopi-app.zip"
fi

TMP=$(mktemp -d)
if ! curl -fsSL --retry 3 -o "$TMP/app.zip" "$URL"; then
    echo "$(date -Is) [$TRACK] download failed: $URL"; rm -rf "$TMP"; exit 0
fi
unzip -q -o "$TMP/app.zip" -d "$TMP/app"
NEW=$(cat "$TMP/app/VERSION")
CUR=$(cat "$APP_DIR/VERSION" 2>/dev/null || echo "none")

if [ "$NEW" = "$CUR" ]; then rm -rf "$TMP"; exit 0; fi # nothing to do

echo "$(date -Is) [$TRACK] deploying $CUR -> $NEW"
mkdir -p "$APP_DIR"
cp -r "$TMP/app/." "$APP_DIR/" # .env is not in the zip, so it is preserved
cd "$APP_DIR" && npm install --omit=dev --no-audit --no-fund
systemctl restart bitopi-app
echo "$(date -Is) [$TRACK] deployed $NEW (commit $(cat "$APP_DIR/COMMIT" 2>/dev/null))"
rm -rf "$TMP"
```

ci-cd.yml

```
# =====
# Bitopi 3-Tier App - CI/CD pipeline (GitHub Actions)
#
# SOURCE : a push to main that touches app/ (or a manual run)
# BUILD  : install dependencies
# TEST   : npm test (Node's built-in test runner against /health, /version, /)
# ARTIFACT: bitopi-app.zip (uploaded as a workflow artifact AND attached
#                           to a GitHub Release tagged with the package.json version)
# DEPLOY  : pull-based - every Auto Scaling instance runs a systemd timer
#           that fetches the newest *stable* release every 60 seconds and
#           restarts the app when the version changes. New instances do
#           the same at boot. No IAM roles or AWS credentials needed.
#
# release_type:
#   stable -> becomes GitHub "latest" -> the BLUE fleet auto-deploys it
#   candidate -> published as pre-release -> ignored by "latest"; only an
#               environment pinned to that tag (GREEN) picks it up
# =====
name: CI/CD - build, test, release

on:
  push:
    branches: [ main ]
    paths: [ 'app/**', '.github/workflows/ci-cd.yml' ]
  workflow_dispatch:
    inputs:
      release_type:
        description: 'stable = Blue auto-deploys · candidate = pre-release for Green'
        type: choice
        options: [ stable, candidate ]
        default: stable

permissions:
  contents: write # needed to create the GitHub Release

jobs:
  build-test-release:
    name: Build → Test → Package → Release
    runs-on: ubuntu-latest
    defaults:
      run:
        working-directory: app

    steps:
      # ----- SOURCE -----
      - name: Checkout source
        uses: actions/checkout@v4

      - name: Set up Node.js 20
        uses: actions/setup-node@v4
        with:
          node-version: 20

      # ----- BUILD -----
      - name: Install dependencies
        run: npm install --no-audit --no-fund

      # ----- TEST -----
      - name: Run application tests
        run: npm test

      # ----- ARTIFACT -----
      - name: Read version from package.json
        id: ver
        run: echo "version=$(node -p "require('./package.json').version")" >> "$GITHUB_OUTPUT"
```

```

- name: Package deployable artifact (bitopi-app.zip)
  run: |
    echo "${{ steps.ver.outputs.version }}" > VERSION
    echo "${{ github.sha }}" > COMMIT
    zip -r ../bitopi-app.zip app.js package.json VERSION COMMIT
    unzip -l ../bitopi-app.zip

- name: Upload artifact to this workflow run
  uses: actions/upload-artifact@v4
  with:
    name: bitopi-app-v${{ steps.ver.outputs.version }}
    path: bitopi-app.zip

# ----- RELEASE (what the servers pull from) -----
- name: Publish GitHub Release v${{ steps.ver.outputs.version }}
  working-directory: ${github.workspace}
  env:
    GH_TOKEN: ${github.token}
    TAG: v${{ steps.ver.outputs.version }}
    TYPE: ${inputs.release_type || 'stable'}
  run: |
    set -e
    PRE=""; [ "$TYPE" = "candidate" ] && PRE="--prerelease"
    NOTES="Built from ${GITHUB_SHA} by GitHub Actions run #${GITHUB_RUN_NUMBER} (${TYPE})."
    if gh release view "$TAG" >/dev/null 2>&1; then
      echo "Release $TAG exists - replacing its asset"
      gh release upload "$TAG" bitopi-app.zip --clobber
    else
      gh release create "$TAG" bitopi-app.zip --title "$TAG" --notes "$NOTES" $PRE
    fi
    echo "::notice title=Released::$TAG ($TYPE) ->
https://github.com/${GITHUB_REPOSITORY}/releases/tag/$TAG"

```

app.js

```
// =====
// Bitopi 3-Tier Application - Backend (Application Tier)
// v1.2.0 - Blue/Green candidate release (GitHub Actions ->
// GitHub Release -> pull-based deployer on every Auto Scaling instance)
//
// GET /health shallow liveness check used by the ALB target group
// GET /version which build is running (used to verify deployments)
// GET / shows which server replied, its environment colour,
// and live data from the MySQL database (Tier 3)
// =====
const express = require('express');
const os = require('os');
const mysql = require('mysql2/promise');

const app = express();
const PORT = process.env.PORT || 3000;
const INSTANCE_ID = process.env.INSTANCE_ID || os.hostname();
const AZ = process.env.AZ || 'unknown';
const ENV_COLOR = (process.env.ENV_COLOR || 'blue').toLowerCase(); // blue | green
const APP_VERSION = require('./package.json').version;

const dbConfig = {
  host: process.env.DB_HOST, user: process.env.DB_USER,
  password: process.env.DB_PASS, database: process.env.DB_NAME,
  port: Number(process.env.DB_PORT) || 3306, connectTimeout: 5000,
};

// ---- /health : shallow on purpose (app alive, NOT a DB check) so a brief
// DB blip never makes the whole fleet unhealthy at once ----
app.get('/health', (req, res) => {
  res.status(200).json({ status: 'healthy', version: APP_VERSION, env: ENV_COLOR,
    instance: INSTANCE_ID, az: AZ });
});

// ---- /version : the one-line answer to "what is deployed right now?" ----
app.get('/version', (req, res) => {
  res.status(200).json({ version: APP_VERSION, env: ENV_COLOR, instance: INSTANCE_ID });
});

// ---- / : proves the full 3-tier round trip ----
app.get('/', async (req, res) => {
  let dbStatus = 'DOWN', dbInfo = '', visitCount = null, dbError = '';
  try {
    const conn = await mysql.createConnection(dbConfig);
    await conn.execute('INSERT INTO visits (hostname) VALUES (?)', [os.hostname()]);
    const [infoRows] = await conn.query('SELECT info FROM app_info WHERE id=1');
    const [cntRows] = await conn.query('SELECT COUNT(*) AS c FROM visits');
    dbInfo = infoRows.length ? infoRows[0].info : '(no info row)';
    visitCount = cntRows[0].c; dbStatus = 'CONNECTED'; await conn.end();
  } catch (e) { dbError = e.message; }

  const accent = ENV_COLOR === 'green' ? '#34d399' : '#38bdf8';
  res.status(200).send(`<!doctype html>
<html><head><meta charset="utf-8"><title>Bitopi 3-Tier App ${APP_VERSION}</title>
<style>
body{font-family:'Segoe UI',Arial,sans-serif;margin:0;background:#0f172a;color:#e2e8f0}
.card{max-width:680px;margin:40px auto;background:#1e293b;border-radius:14px;padding:28px
32px;box-shadow:0 10px 30px rgba(0,0,0,.4);border-top:6px solid ${accent}}
h1{margin:0 0 6px;color:${accent}}.k{color:#94a3b8}.v{color:#f8fafc;font-weight:600}
.row{display:flex;justify-content:space-between;padding:10px 0;border-bottom:1px solid #334155}
.ok{color:#34d399;font-weight:700}.bad{color:#f87171;font-weight:700}
.badge{display:inline-block;background:${accent};color:#08131f;padding:2px 10px;border-
radius:20px;font-size:12px;font-weight:700;margin:0 6px 14px 0}
.ver{font-size:26px;font-weight:800;color:${accent}}
</style></head><body><div class="card">
<h1>Bitopi 3-Tier Application</h1>
```

```

<span class="badge">BACKEND TIER &middot; AUTO SCALING</span><span
class="badge">${ENV_COLOR.toUpperCase()} ENVIRONMENT</span>
<div class="row"><span class="k">Application version</span><span class="ver">v${APP_VERSION}
</span></div>
<div class="row"><span class="k">Environment</span><span class="v">${ENV_COLOR}</span></div>
<div class="row"><span class="k">Served by instance</span><span class="v">${INSTANCE_ID}</span>
</div>
<div class="row"><span class="k">Availability Zone</span><span class="v">${AZ}</span></div>
<div class="row"><span class="k">Hostname</span><span class="v">${os.hostname()}</span></div>
<div class="row"><span class="k">Database (Tier 3)</span><span
class="${dbStatus==='CONNECTED'? 'ok': 'bad'}">${dbStatus}</span></div>
${dbStatus==='CONNECTED'
? `<div class="row"><span class="k">DB message</span><span class="v">${dbInfo}</span></div>
  <div class="row"><span class="k">Total visits recorded</span><span class="v">${visitCount}
</span></div>`
: `<div class="row"><span class="k">DB error</span><span class="bad">${dbError}</span></div>`}
<p style="color:#64748b;margin-top:18px;font-size:13px">Deployed by GitHub Actions &arr; GitHub
Release &arr; pull-based deployer on this server. Refresh to be routed to another instance.</p>
<p style="color:#fbbf24;font-size:13px;font-weight:600">What's new in v${APP_VERSION}:
${ENV_COLOR==='green'? 'Blue/Green candidate &mdash; this version runs ONLY on the GREEN
environment until the ALB listener is switched.' : 'this line reached the server through
automatic pull-based deployment &mdash; no AWS console action was taken.'}</p>
</div></body></html>`);
});

// Only start listening when run directly (tests import the app instead)
if (require.main === module) {
  app.listen(PORT, () => console.log(`Bitopi app v${APP_VERSION} (${ENV_COLOR}) listening on port
${PORT}`));
}
module.exports = app;

```

app-server-userdata-v3-green.sh

Identical to `app-server-userdata-v2-blue.sh` above except for three values: `ENV_COLOR=green`, `DEPLOY_TRACK=v1.2.0` and the log line. Full file: `scripts/app-server-userdata-v3-green.sh`.

terraform/main.tf

```
# =====
# Bitopi 3-Tier assignment - Terraform "is it still there?" audit
#
# This configuration is READ-ONLY. It contains ZERO `resource` blocks,
# so `terraform apply` can never create, change or delete anything in
# AWS. It only *looks up* the resources that were built by hand in
# Part 1 / Part 2 and reports, for each one, PRESENT or MISSING.
#
# Files:
#   main.tf       - provider / version settings + the names we look for
#   audit.tf      - the lookups, the PRESENT/MISSING checks, the report
#   README.md     - how to run it on Windows, step by step
# =====

terraform {
  required_version = ">= 1.5.0" # `check` blocks need Terraform 1.5+

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 5.40"
    }
  }
}

provider "aws" {
  region = var.region
  # Credentials are NOT written here. Terraform reads them from
  # environment variables (AWS_ACCESS_KEY_ID / AWS_SECRET_ACCESS_KEY /
  # AWS_SESSION_TOKEN) or from `aws configure` - see README.md.
}

# ---- names used when the stack was built (change only if you renamed) ----
variable "region" {
  description = "Region the stack was built in"
  type        = string
  default     = "ap-south-1"
}

variable "prefix" {
  description = "Name prefix used for every resource of the assignment"
  type        = string
  default     = "bitopi"
}

locals {
  vpc_name      = "${var.prefix}-vpc"
  subnet_glob   = "${var.prefix}-subnet-*" # bitopi-subnet-public1-ap-south-1a, ...
  sg_names      = ["${var.prefix}-alb-sg", "${var.prefix}-backend-sg", "${var.prefix}-db-sg"]
  key_pair_name = "${var.prefix}-key"
  db_server_name = "${var.prefix}-db-server"
  launch_template = "${var.prefix}-app-lt"
  tg_blue       = "${var.prefix}-app-tg"
  tg_green      = "${var.prefix}-app-tg-green"
  alb_name      = "${var.prefix}-alb"
  asg_blue      = "${var.prefix}-app-asg"
  asg_green     = "${var.prefix}-app-asg-green"
}
```

terraform/audit.tf

```
# =====
# audit.tf - look up every hand-built resource and report on it
#
# Two kinds of lookups are used on purpose:
#
# 1. "plural" data sources (aws_vpcs, aws_subnets, aws_security_groups,
#    aws_instances, aws_lbs, aws_autoscaling_groups). These NEVER fail:
#    when nothing matches they simply return an empty list. That lets
#    us print PRESENT / MISSING in the final report and also print
#    useful values (DB private IP, ALB DNS name, ASG sizes ...).
#
# 2. `check` blocks with a "scoped" data source inside them. Some
#    resources only have a singular data source (key pair, launch
#    template, target group) which would normally ABORT the whole run
#    when the thing does not exist. Inside a check block Terraform
#    turns that abort into a *warning* instead, so the run continues
#    and the warning itself is the MISSING report.
# =====

# -----
# 1. Network - VPC, subnets, security groups
# -----
data "aws_vpcs" "vpc" {
  filter {
    name     = "tag:Name"
    values   = [local.vpc_name]
  }
}

data "aws_subnets" "subnets" {
  filter {
    name     = "tag:Name"
    values   = [local.subnet_glob]
  }
}

data "aws_subnet" "each" {
  for_each = toset(data.aws_subnets.subnets.ids)
  id       = each.value
}

data "aws_security_groups" "sgs" {
  filter {
    name     = "group-name"
    values   = local.sg_names
  }
}

data "aws_security_group" "each" {
  for_each = toset(data.aws_security_groups.sgs.ids)
  id       = each.value
}

# -----
# 2. Database tier - the MySQL (MariaDB) EC2 server
# -----
data "aws_instances" "db" {
  filter {
    name     = "tag:Name"
    values   = [local.db_server_name]
  }
  instance_state_names = ["pending", "running", "stopping", "stopped"]
}

# -----
# 3. Load balancer + listener (list all ELBv2 LBs, then pick ours)
```

```

# -----
data "aws_lbs" "all" {}

data "aws_lb" "each" {
  for_each = data.aws_lbs.all.arns
  arn      = each.value
}

locals {
  alb = one([for lb in data.aws_lb.each : lb if lb.name == local.alb_name])
}

data "aws_lb_listener" "http80" {
  count          = local.alb == null ? 0 : 1
  load_balancer_arn = local.alb.arn
  port          = 80
}

# -----
# 4. Auto Scaling Groups (blue = Part 1/2, green = Blue/Green step)
# -----
data "aws_autoscaling_groups" "bitopi" {
  names = [local.asg_blue, local.asg_green]
}

data "aws_autoscaling_group" "each" {
  for_each = toset(data.aws_autoscaling_groups.bitopi.names)
  name     = each.value
}

# -----
# 5. Things that only have a singular lookup -> check blocks
# (a failure here prints a WARNING, never an error)
# -----
check "key_pair" {
  data "aws_key_pair" "this" {
    key_name = local.key_pair_name
  }
  assert {
    condition     = data.aws_key_pair.this.key_name == local.key_pair_name
    error_message = "MISSING: key pair ${local.key_pair_name} (Part 1 - Step 2)"
  }
}

check "launch_template" {
  data "aws_launch_template" "this" {
    name = local.launch_template
  }
  assert {
    condition     = data.aws_launch_template.this.latest_version >= 2
    error_message = "Launch template ${local.launch_template} exists but only has version
${data.aws_launch_template.this.latest_version}. Part 2 needs v2 (pull-based deployer) - see
docs/part2-cicd.md Step 2."
  }
}

check "target_group_blue" {
  data "aws_lb_target_group" "this" {
    name = local.tg_blue
  }
  assert {
    condition     = data.aws_lb_target_group.this.port == 3000 &&
data.aws_lb_target_group.this.health_check[0].path == "/health"
    error_message = "Target group ${local.tg_blue} exists but is not HTTP:3000 with health check
/health."
  }
}

```

```

}

check "target_group_green_optional" {
  data "aws_lb_target_group" "green" {
    name = local.tg_green
  }
  assert {
    condition      = data.aws_lb_target_group.green.port == 3000
    error_message = "Green target group ${local.tg_green} exists but is not on port 3000."
  }
}

# -----
# 6. Assertions on the plural lookups (so they also show as warnings)
# -----
check "vpc" {
  assert {
    condition      = length(data.aws_vpcs.vpc.ids) == 1
    error_message = "MISSING: VPC ${local.vpc_name} (Part 1 - Step 1). Found
${length(data.aws_vpcs.vpc.ids)}."
  }
}

check "subnets" {
  assert {
    condition      = length(data.aws_subnets.subnets.ids) == 4
    error_message = "Expected 4 subnets named ${local.subnet_glob}, found
${length(data.aws_subnets.subnets.ids)} (Part 1 - Step 1).\"
  }
}

check "security_groups" {
  assert {
    condition      = length(data.aws_security_groups.sgs.ids) == 3
    error_message = "Expected 3 security groups (${join(", ", local.sg_names)}), found
${length(data.aws_security_groups.sgs.ids)} (Part 1 - Step 2).\"
  }
}

check "db_server" {
  assert {
    condition      = length(data.aws_instances.db.ids) >= 1
    error_message = "MISSING: EC2 instance ${local.db_server_name} (Part 1 - Step 3). The app's
DB_HOST must be rebuilt.\"
  }
}

check "alb" {
  assert {
    condition      = local.alb != null
    error_message = "MISSING: Application Load Balancer ${local.alb_name} (Part 1 - Step 5).\"
  }
}

check "asg_blue" {
  assert {
    condition      = contains(data.aws_autoscaling_groups.bitopi.names, local.asg_blue)
    error_message = "MISSING: Auto Scaling Group ${local.asg_blue} (Part 1 - Step 6).\"
  }
}

# -----
# 7. The report
# -----
locals {
  ok = {

```

```

vpc      = length(data.aws_vpcs.vpc.ids) == 1
subnets = length(data.aws_subnets.subnets.ids) == 4
sgs      = length(data.aws_security_groups.sgs.ids) == 3
db        = length(data.aws_instances.db.ids) >= 1
alb       = local.alb != null
asg_blue  = contains(data.aws_autoscaling_groups.bitopi.names, local.asg_blue)
asg_green = contains(data.aws_autoscaling_groups.bitopi.names, local.asg_green)
}
mark = { for k, v in local.ok : k => (v ? "PRESENT" : "MISSING") }

blue_asg = try(data.aws_autoscaling_group.each[local.asg_blue], null)
green_asg = try(data.aws_autoscaling_group.each[local.asg_green], null)
listener = try(data.aws_lb_listener.http80[0], null)
# arn:aws:elasticloadbalancing:::targetgroup/<NAME>/<hash> -> <NAME>
live_tg = local.listener == null ? "-" : try(regex("targetgroup/([^\s]+)/",
local.listener.default_action[0].target_group_arn)[0], "?")

db_ip    = length(data.aws_instances.db.private_ips) > 0 ? data.aws_instances.db.private_ips[0]
: null
alb_dns = local.alb == null ? null : local.alb.dns_name

report = <<-EOT

===== BITOPI 3-TIER AUDIT (${var.region}) =====
VPC ${local.vpc_name} ..... ${local.mark.vpc}
4 subnets ${local.subnet_glob} ..... ${local.mark.subnets} (found
${length(data.aws_subnets.subnets.ids)})
3 security groups ..... ${local.mark.sgs} (found
${length(data.aws_security_groups.sgs.ids)}: ${join(", ", [for sg in data.aws_security_group.each
: sg.name])})
DB server ${local.db_server_name} ..... ${local.mark.db}   ${local.db_ip == null ?
"" : "private IP ${local.db_ip}"}
ALB ${local.alb_name} ..... ${local.mark.alb}   ${local.alb_dns ==
null ? "" : local.alb_dns}
listener :80 currently forwards to ..... ${local.live_tg}
ASG ${local.asg_blue} (blue) ..... ${local.mark.asg_blue}   ${local.blue_asg ==
null ? "" : "min ${local.blue_asg.min_size} / desired ${local.blue_asg.desired_capacity} / max
${local.blue_asg.max_size}, LT ${try(local.blue_asg.launch_template[0].name, "?")}}
v${try(local.blue_asg.launch_template[0].version, "?")}"
ASG ${local.asg_green} (green) ..... ${local.mark.asg_green} (only needed for the
Blue/Green step)
key pair / launch template / target groups : see the WARNINGS above
                                         (no warning = PRESENT)

=====

Everything marked MISSING (or named in a warning) must be rebuilt by
hand in the console - follow docs/part1-auto-scaling.md for Part 1 and
docs/part2-cicd.md section 7 for the Blue/Green steps.
EOT
}

output "audit_summary" {
  description = "Human-readable PRESENT/MISSING report"
  value       = local.report
}

output "db_private_ip" {
  description = "Use this as DB_HOST in the launch-template user-data (null = DB server missing)"
  value       = local.db_ip
}

output "alb_url" {
  description = "Open this in the browser (null = ALB missing)"
  value       = local.alb_dns == null ? null : "http://${local.alb_dns}/"
}

```

```

output "public_subnet_ids" {
  description = "Subnets to pick for the ALB and the Auto Scaling Groups"
  value       = { for s in data.aws_subnet.each : s.tags["Name"] => s.id if can(regex("public",
s.tags["Name"])) }
}

output "security_group_ids" {
  value = { for sg in data.aws_security_group.each : sg.name => sg.id }
}

output "blue_asg" {
  value = local.blue_asg == null ? null : {
    min                = local.blue_asg.min_size
    desired             = local.blue_asg.desired_capacity
    max                = local.blue_asg.max_size
    launch_template_version = try(local.blue_asg.launch_template[0].version, null)
    health_check_type    = local.blue_asg.health_check_type
    target_groups        = local.blue_asg.target_group_arns
  }
}

output "green_asg" {
  value = local.green_asg == null ? null : {
    min                = local.green_asg.min_size
    desired             = local.green_asg.desired_capacity
    max                = local.green_asg.max_size
    launch_template_version = try(local.green_asg.launch_template[0].version, null)
  }
}

output "listener_forwards_to" {
  description = "Which target group gets production traffic right now (blue = bitopi-app-tg,
green = bitopi-app-tg-green)"
  value       = local.live_tg
}

```